

# SQL: Schemas, Queries, Updates, Views

---

## SQL

1/198

SQL = **Structured Query Language** (sometimes called "sequel").

SQL is an ANSI/ISO standard language for querying and manipulating relational DBMSs.

Designed to be a "human readable" language comprising:

- data definition facilities
  - database modification operations
  - database query operations, including:
    - relational algebra, set operations, aggregation, grouping, ...
- 

## ... SQL

2/198

SQL was developed at IBM (San Jose Lab) during the 1970's, and standardised in 1986.

DBMSs typically implement the SQL2 standard (aka SQL-92).

Unfortunately, they also:

- implement a (large) subset of the standard
- extend the standard in various "useful" ways

SQL (in some form) looks likely to survive in the next generation of database systems.

In these slides, we try to use only *standard* (portable) SQL2.

---

## ... SQL

3/198

Since SQL2, there have been three new proposed standards:

SQL:1999 added e.g.

- boolean and BLOB types, arrays/rows, ...
- procedures programming constructs, triggers
- recursive queries
- OO-like objects, inheritance, ...

SQL:2003 ...

- standardised some SQL:1999 extensions
- added a standard for meta-data (catalogues)
- standardised stored procedures (SQL/PSM)
- added a new MERGE statement ("upsert")
- defined interfaces to C, Java, XML, object systems, ...

SQL:2008 added additional support for XML.

---

## ... SQL

4/198

Major DBMSs (Oracle, DB2, SQLServer, PostgreSQL MySQL):

- implement most/all of SQL2
- implement much of SQL:1999
- implement some of SQL:2003
- omit difficult-to-implement features e.g. assertions

PostgreSQL ...

- implements almost all of SQL2 (see documentation)
- does not implement: recursive queries, assertions

- provides non-standard mechanisms for: updatable views
- currently has PLpgSQL, will also have SQL/PSM soon

## ... SQL

5/198

SQL provides high-level, declarative access to data.

However, SQL is not a Turing-complete programming language.

Applications typically embed evaluation of SQL queries into PL's:

- Java and the JDBC API
- PHP/Perl/Tcl and their various DBMS bindings
- RDBMS-specific programming languages  
(e.g. Oracle's PL/SQL, PostgreSQL's PLpgSQL)
- C and low-level library interfaces to DBMS engine  
(e.g. Oracle's OCI, PostgreSQL's libpq)

## ... SQL

6/198

SQL's query sub-language is based on *relational algebra*.

Relational algebra:

- formal language of expressions mapping tables→tables
- comprising three basic operations ...
  - *select*: filter table rows via a condition on attributes
  - *project*: filter table columns by name
  - *join*: combines two tables via a condition
- along with set operations (union, intersection, difference)
- and a variety of aggregates (including min(), max(), count(), etc)

## ... SQL

7/198

Example relational algebra operations:

| R                 | <table><tr><th>t</th><th>u</th><th>v</th></tr><tr><td>1</td><td>'a'</td><td>2.5</td></tr><tr><td>1</td><td>'b'</td><td>1.7</td></tr><tr><td>2</td><td>'a'</td><td>3.0</td></tr><tr><td>2</td><td>'b'</td><td>2.4</td></tr></table>   | t   | u   | v | 1 | 'a' | 2.5 | 1   | 'b' | 1.7 | 2 | 'a' | 3.0 | 2   | 'b' | 2.4 | S | <table><tr><th>x</th><th>y</th></tr><tr><td>'h'</td><td>1</td></tr><tr><td>'j'</td><td>2</td></tr><tr><td>'k'</td><td>3</td></tr></table> | x   | y   | 'h' | 1 | 'j' | 2   | 'k' | 3 |  |  |
|-------------------|--|-----|-----|---|---|-----|-----|-----|-----|-----|---|-----|-----|-----|-----|-----|---|---|-----|-----|-----|---|-----|-----|-----|---|--|--|
| t                 | u  | v   |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'a'  | 2.5 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'b'  | 1.7 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 2                 | 'a'  | 3.0 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 2                 | 'b'  | 2.4 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| x                 | y  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 'h'               | 1  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 'j'               | 2  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 'k'               | 3  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| select [t=1] (R)  | <table><tr><th>t</th><th>u</th><th>v</th></tr><tr><td>1</td><td>'a'</td><td>2.5</td></tr><tr><td>1</td><td>'b'</td><td>1.7</td></tr></table>   | t   | u   | v | 1 | 'a' | 2.5 | 1   | 'b' | 1.7 |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| t                 | u  | v   |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'a'  | 2.5 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'b'  | 1.7 |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| project [t] (R)   | <table><tr><th>t</th></tr><tr><td>1</td></tr><tr><td>2</td></tr></table>   | t   | 1   | 2 |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| t                 |  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 |  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 2                 |  |     |     |   |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| join [t=y] (R, S) | <table><tr><th>t</th><th>u</th><th>v</th><th>x</th><th>y</th></tr><tr><td>1</td><td>'a'</td><td>2.5</td><td>'h'</td><td>1</td></tr><tr><td>1</td><td>'b'</td><td>1.7</td><td>'h'</td><td>1</td></tr><tr><td>2</td><td>'a'</td><td>3.0</td><td>'j'</td><td>2</td></tr><tr><td>2</td><td>'b'</td><td>2.4</td><td>'j'</td><td>2</td></tr></table> | t   | u   | v | x | y   | 1   | 'a' | 2.5 | 'h' | 1 | 1   | 'b' | 1.7 | 'h' | 1   | 2 | 'a'   | 3.0 | 'j' | 2   | 2 | 'b' | 2.4 | 'j' | 2 |  |  |
| t                 | u  | v   | x   | y |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'a'  | 2.5 | 'h' | 1 |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 1                 | 'b'  | 1.7 | 'h' | 1 |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 2                 | 'a'  | 3.0 | 'j' | 2 |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |
| 2                 | 'b'  | 2.4 | 'j' | 2 |   |     |     |     |     |     |   |     |     |     |     |     |   |   |     |     |     |   |     |     |     |   |  |  |

## Example Databases

8/198

In order to demonstrate aspects of SQL, we use two databases:


- bank: customers, accounts, branches, ...
- beer: beers, bars, drinkers, ...

These databases are available for you to play with.

## Example Database #1

9/198

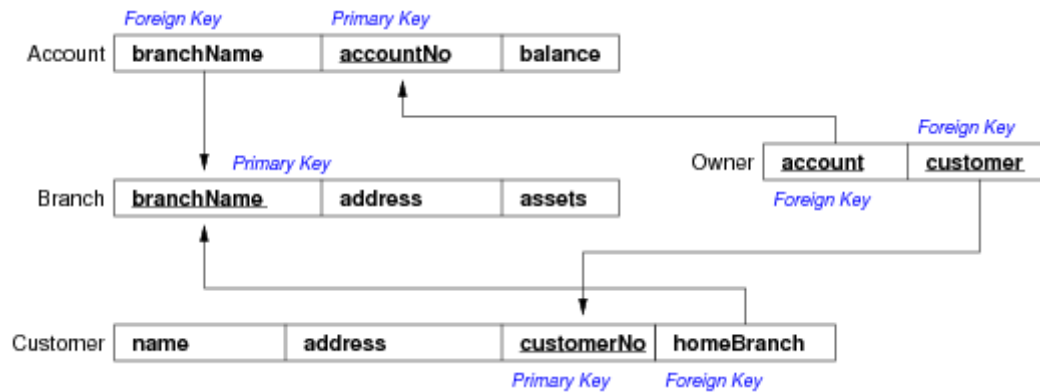
ER design for a simple banking application:



## ... Example Database #1

10/198

Relational schema corresponding to the ER design:



## ... Example Database #1

11/198

We will use the following instance of this schema:

**Branch** relation/table instance:

| branchName | address        | assets |
|------------|----------------|--------|
| Clovelly   | Clovelly Rd.   | 1000   |
| Coogee     | Coogee Bay Rd. | 40000  |
| Maroubra   | Anzac Pde.     | 17000  |
| Randwick   | Alison Rd.     | 20000  |
| UNSW       | near Library   | 3000   |

**Customer** relation/table instance:

| name   | address        | customerNo | homebranch |
|--------|----------------|------------|------------|
| Adam   | Belmore Rd.    | 12345      | Randwick   |
| Bob    | Rainbow St.    | 32451      | Coogee     |
| Chuck  | Clovelly Rd.   | 76543      | Clovelly   |
| David  | Anzac Pde.     | 82199      | UNSW       |
| George | Anzac Pde.     | 81244      | Maroubra   |
| Graham | Malabar Rd.    | 92754      | Maroubra   |
| Greg   | Coogee Bay Rd. | 22735      | Coogee     |
| Jack   | High St.       | 12666      | Randwick   |

## ... Example Database #1

12/198

**Account** relation/table instance:

| branchName | accountNo | balance |
|------------|-----------|---------|
| -----      | -----     | -----   |

|          |       |       |
|----------|-------|-------|
| UNSW     | U-245 | 1000  |
| UNSW     | U-291 | 2000  |
| Randwick | R-245 | 20000 |
| Coogee   | C-123 | 15000 |
| Coogee   | C-124 | 25000 |
| Clovelly | Y-123 | 1000  |
| Maroubra | M-222 | 5000  |
| Maroubra | M-225 | 12000 |

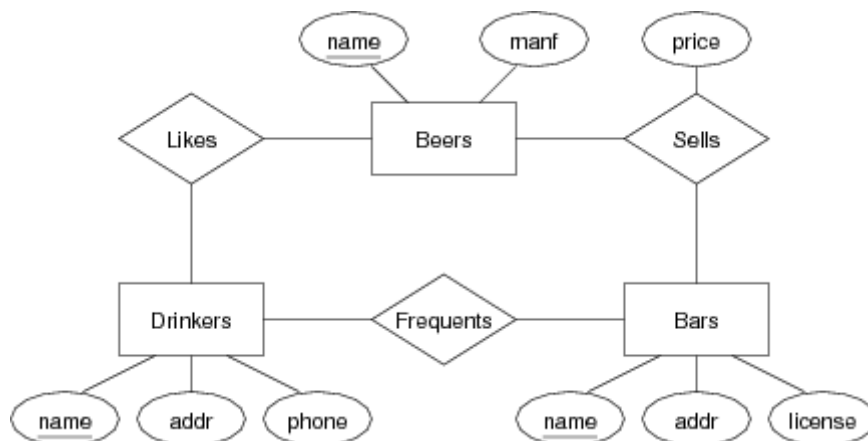
**Owner** relation/table instance:

| account | customer |
|---------|----------|
| U-245   | 12345    |
| U-291   | 12345    |
| U-291   | 12666    |
| R-245   | 12666    |
| C-123   | 32451    |
| C-124   | 22735    |
| Y-123   | 76543    |
| M-222   | 92754    |
| M-225   | 12345    |

## Example Database #2

13/198

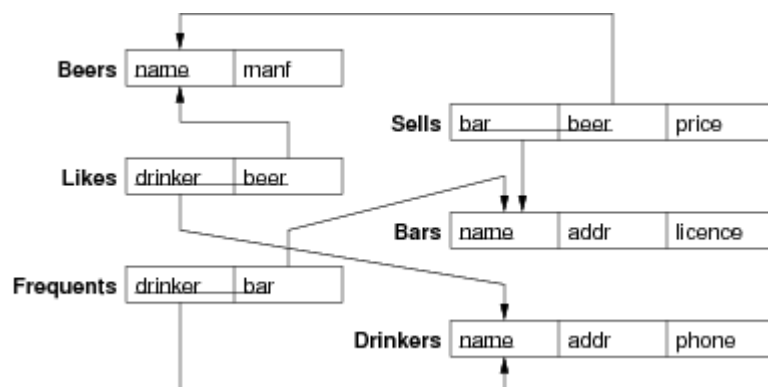
ER design for beers/bars/drinkers database:



## ... Example Database #2

14/198

Relational schema corresponding to the ER design:



## ... Example Database #2

15/198

We will use the following instance of this schema:

**Bars** relation/table instance:

| name | addr | license |
|------|------|---------|
|------|------|---------|

|                  |           |        |
|------------------|-----------|--------|
| Australia Hotel  | The Rocks | 123456 |
| Coogee Bay Hotel | Coogee    | 966500 |
| Lord Nelson      | The Rocks | 123888 |
| Marble Bar       | Sydney    | 122123 |
| Regent Hotel     | Kingsford | 987654 |
| Royal Hotel      | Randwick  | 938500 |

**Drinkers** relation/table instance:

| name   | addr     | phone     |
|--------|----------|-----------|
| Adam   | Randwick | 9385-4444 |
| Gernot | Newtown  | 9415-3378 |
| John   | Clovelly | 9665-1234 |
| Justin | Mosman   | 9845-4321 |

## ... Example Database #2

16/198

**Beers** relation/table instance:

| name                | manf          |
|---------------------|---------------|
| 80/-                | Caledonian    |
| Bigfoot Barley Wine | Sierra Nevada |
| Burraborang Bock    | George IV Inn |
| Crown Lager         | Carlton       |
| Fosters Lager       | Carlton       |
| Invalid Stout       | Carlton       |
| Melbourne Bitter    | Carlton       |
| New                 | Toohey's      |
| Old                 | Toohey's      |
| Old Admiral         | Lord Nelson   |
| Pale Ale            | Sierra Nevada |
| Premium Lager       | Cascade       |
| Red                 | Toohey's      |
| Sheaf Stout         | Toohey's      |
| Sparkling Ale       | Cooper's      |
| Stout               | Cooper's      |
| Three Sheets        | Lord Nelson   |
| Victoria Bitter     | Carlton       |

## ... Example Database #2

17/198

**Frequents** relation/table instance:

| drinker | bar              |
|---------|------------------|
| Adam    | Coogee Bay Hotel |
| Gernot  | Lord Nelson      |
| John    | Coogee Bay Hotel |
| John    | Lord Nelson      |
| John    | Australia Hotel  |
| Justin  | Regent Hotel     |
| Justin  | Marble Bar       |

## ... Example Database #2

18/198

**Likes** relation/table instance:

| drinker | beer                |
|---------|---------------------|
| Adam    | Crown Lager         |
| Adam    | Fosters Lager       |
| Adam    | New                 |
| Gernot  | Premium Lager       |
| Gernot  | Sparkling Ale       |
| John    | 80/-                |
| John    | Bigfoot Barley Wine |

|        |  |                 |
|--------|--|-----------------|
| John   |  | Pale Ale        |
| John   |  | Three Sheets    |
| Justin |  | Sparkling Ale   |
| Justin |  | Victoria Bitter |

## ... Example Database #2

19/198

**Sells** relation/table instance:

| bar              | beer             | price |
|------------------|------------------|-------|
| Australia Hotel  | Burraborang Bock | 3.50  |
| Coogee Bay Hotel | New              | 2.25  |
| Coogee Bay Hotel | Old              | 2.50  |
| Coogee Bay Hotel | Sparkling Ale    | 2.80  |
| Coogee Bay Hotel | Victoria Bitter  | 2.30  |
| Lord Nelson      | Three Sheets     | 3.75  |
| Lord Nelson      | Old Admiral      | 3.75  |
| Marble Bar       | New              | 2.80  |
| Marble Bar       | Old              | 2.80  |
| Marble Bar       | Victoria Bitter  | 2.80  |
| Regent Hotel     | New              | 2.20  |
| Regent Hotel     | Victoria Bitter  | 2.20  |
| Royal Hotel      | New              | 2.30  |
| Royal Hotel      | Old              | 2.30  |
| Royal Hotel      | Victoria Bitter  | 2.30  |

## SQL Syntax

20/198

SQL definitions, queries and statements are composed of:

- *comments* ... `--` comments to end of line
- *identifiers* ... similar to regular programming languages
- *keywords* ... a large set (e.g. CREATE, SELECT, TABLE)
- *data types* ... a small set (e.g. integer, varchar, date)
- *operators* ... similar to regular programming languages
- *constants* ... similar to regular programming languages

*Similar* means "often the same, but not always ..."

- 'John', 'blue', 'it's' are *strings*
- "Students", "Really Silly!" are *identifiers*

## ... SQL Syntax

21/198

While SQL identifiers and keywords are case-insensitive, we generally:

- write keywords in upper case (until it becomes annoying)  
e.g. SELECT, FROM, WHERE, CREATE, ...
- write relation names with an initial upper-case letter  
e.g. Customers, Students, Owns, EnrolledIn
- write attribute names in all lower-case  
e.g. id, name, partNumber, isActive

We follow the above conventions when writing programs.

We ignore the above conventions when typing in lectures.

## SQL Keywords

22/198

A categorised list of frequently-used SQL92 keywords:

| Querying | Defining Data | Changing Data |
|----------|---------------|---------------|
| SELECT   | CREATE        | INSERT        |
| FROM     | TABLE         | INTO          |
| WHERE    | INTEGER       | VALUES        |

|          |            |        |
|----------|------------|--------|
| GROUP BY | REAL       | UPDATE |
| HAVING   | VARCHAR    | SET    |
| ORDER BY | CHAR       | DELETE |
| DESC     | KEY        | DROP   |
| EXISTS   | PRIMARY    | ALTER  |
| IS NULL  | FOREIGN    |        |
| NOT NULL | REFERENCES |        |
| IN       | CONSTRAINT |        |
| DISTINCT | CHECK      |        |
| AS       |            |        |

There are 225 reserved words in SQL92 ... not a small language.

## ... SQL Keywords

23/198

A list of PostgreSQL's SQL keywords:

|              |            |         |              |
|--------------|------------|---------|--------------|
| ALL          | DEFERRABLE | IS      | OVERLAPS     |
| ANALYSE      | DESC       | ISNULL  | PRIMARY      |
| ANALYZE      | DISTINCT   | JOIN    | PUBLIC       |
| AND          | DO         | LEADING | REFERENCES   |
| ANY          | ELSE       | LEFT    | RIGHT        |
| AS           | END        | LIKE    | SELECT       |
| ASC          | EXCEPT     | LIMIT   | SESSION_USER |
| BETWEEN      | FALSE      | NATURAL | SOME         |
| BINARY       | FOR        | NEW     | TABLE        |
| BOTH         | FOREIGN    | NOT     | THEN         |
| CASE         | FREEZE     | NOTNULL | TO           |
| CAST         | FROM       | NULL    | TRAILING     |
| CHECK        | FULL       | OFF     | TRUE         |
| COLLATE      | GROUP      | OFFSET  | UNION        |
| COLUMN       | HAVING     | OLD     | UNIQUE       |
| CONSTRAINT   | ILIKE      | ON      | USER         |
| CROSS        | IN         | ONLY    | USING        |
| CURRENT_DATE | INITIALLY  | OR      | VERBOSE      |
| CURRENT_TIME | INNER      | ORDER   | WHEN         |
| CURRENT_USER | INTERSECT  | OUTER   | WHERE        |
| DEFAULT      | INTO       |         |              |

Note that some SQL92 reserved words are not reserved words in PostgreSQL.

## SQL Identifiers

24/198

Names are used to identify

- database objects such as tables, attributes, views, ...
- meta-objects such as types, functions, constraints, ...

Identifiers in SQL use similar conventions to programming languages i.e. a sequence of alpha-numerics, starting with an alphabetic.

Can create arbitrary identifiers by enclosing in "..."

Example identifiers:

```
employee    student    Courses
last_name   "That's a Great Name!"
```

Oracle SQL also allows unquoted hash (#) and dollar (\$) in identifiers.

## ... SQL Identifiers

25/198

Since SQL does not distinguish case, the following are all treated as being the same identifier:

```
employee    Employee    EmPlOyEe
```

Most RDBMSs will let you give the same name to different kinds of objects (e.g. a table called Beer and an attribute called Beer).

Some common naming conventions:

- name tables representing entities via plural nouns (e.g. Drinkers, TheDrinkers, AllDrinkers, ...)
- name foreign key attributes after the table they refer to (e.g. beer in the Sells relation)

---

## Constants in SQL

26/198

Numeric constants have same syntax as programming languages, e.g.

10      3.14159      2e-5      6.022e23

String constants are written in single quotes, e.g.

'John'      'some text'      '!%#%!\$'      'O''Brien'  
 '''      '[A-Z]{4}\d{4}'      'a VeRy! LoNg String'

PostgreSQL provides extended strings containing \ escapes, e.g.

E'\n'      E'O\'Brien'      E'[A-Z]{4}\\d{4}'      E'John'

Boolean constants: TRUE and FALSE

PostgreSQL also allows 't', 'true', 'yes', 'f', 'false', 'no'

---

## ... Constants in SQL

27/198

Other kinds of constants are typically written as strings.

Dates: '2008-04-13', Times: '13:30:15'

Timestamps: '2004-10-19 10:23:54'

PostgreSQL also recognises: 'January 26 11:05:10 1988 EST'

Time intervals: '10 minutes', '5 days, 6 hours'

PostgreSQL also has IP address, XML, etc. data types.

---

## SQL Data Types

28/198

All attributes in SQL relations are typed (i.e. have domain specified)

SQL supports a small set of useful built-in data types:  
 text string, number (integer,real), date, boolean, binary

Various type conversions are available (e.g. date to string, string to date, integer to real) and applied automatically "where they make sense".

Basic domain (type) checking is performed automatically.

The NULL value is treated as a member of all data types.

No structured data types are available (in SQL2).

---

## ... SQL Data Types

29/198

Various kinds of number types are available:

- INTEGER (or INT), SMALLINT ... 32/16-bit integers
- REAL, DOUBLE PRECISION ... 32/64-bit floating point
- NUMBER(*d*,*p*) ... fixed-point reals (*d* digits, *p* after dec.pt.)

PostgreSQL also provides ...

- serial: auto-generated integer values for primary keys
- currency: fixed-point reals, displayed as strings \$1,000.00



Two string types are available:

- `CHAR(n)` ... uses *n* bytes, left-justified, blank-padded
- `VARCHAR(n)` ... uses 0..*n* bytes, no padding

String types can be coerced by blank-padding or truncation.

```
'abc'::CHAR(2) = 'ab'      'abc'::CHAR(4) = 'abc '
```

PostgreSQL also provides `TEXT` for arbitrary strings

- convenient; no need to worry "how long is a name?"
- efficient (different to some other DBMSs)
- but not part of SQL standard

Dates are simply specially-formatted strings, with a range of operations to implement date semantics.

Format is typically `YYYY-MM-DD`, e.g. `'1998-08-02'`

Accepts other formats (and has format-conversion functions), but beware of two-digit years (year 2000)

Comparison operators implement *before* (`<`) and *after* (`>`).

Subtraction counts number of days between two dates.

Etc. etc. ... consult your local SQL Manual

PostgreSQL also supports several non-standard data types.

- generic text string data i.e. `text`
- arbitrary binary data (BLOBs) i.e. `bytea`
- geometric data types e.g. `point`, `circle`, `polygon`, ...

Also, extends relational model so that a single attribute can contain an array/matrix of values, e.g.

```
CREATE TABLE Employees (  
    empid      integer primary key,  
    name       text,  
    pay_rate   float[]  
);  
INSERT INTO Employees VALUES  
    (1234, 'John', '{35.00,45.00,60.00}');  
SELECT pay_rate[2] FROM Employees ...
```

Tuple and set constants are both written as:

```
( val1, val2, val3, ... )
```

The correct interpretation is worked out from the context.

Examples:

```
INSERT INTO Student(stude#, name, course)  
VALUES (2177364, 'Jack Smith', 'BSc')  
-- tuple literal
```

```
CREATE TABLE Academics (  
    id integer,
```

```
name varchar(40),
job  varchar(10) CHECK
      job IN ('Lecturer', 'Tutor');
-- set literal
```

---

## ... Tuple and Set Literals

34/198

SQL data types provide coarse-grained control over values.

If more fine-grained control over values is needed:

- constraints can express more precise conditions
- new "data types" can be defined

Examples:

```
CREATE DOMAIN PositiveInt AS INTEGER
CHECK (VALUE > 0);
CREATE DOMAIN Colour AS
CHECK (VALUE IN ('red','yellow','green','blue','violet'));
CREATE TABLE T (
  x Colour,
  y PositiveInt,
  z INTEGER CHECK (z BETWEEN 10 AND 20)
);
```

---

## SQL Operators

35/198

Comparison operators are defined on all types:

<   >   <=   >=   =   <>   (or !=)

Boolean operators AND, OR, NOT are also available

Note AND,OR are not "short-circuit" in the same way as C's &&, ||

Most data types also have type-specific operations available

See PostgreSQL Documentation Chapter 8/9 for data types and operators

---

## ... SQL Operators

36/198

**String comparison:**

- $str_1 < str_2$  ... compare using dictionary order
- $str \text{ LIKE } pattern$  ... matches string to pattern

Pattern-matching uses SQL-specific pattern expressions:

- % matches anything (like .\*)
  - \_ matches any single char (like .)
- 

## ... SQL Operators

37/198

Examples (using SQL92 pattern matching):

|                   |                            |
|-------------------|----------------------------|
| Name LIKE 'Ja%'   | Name begins with 'Ja'      |
| Name LIKE '_i%'   | Name has 'i' as 2nd letter |
| Name LIKE '%o%o%' | Name contains two 'o's     |
| Name LIKE '%ith'  | Name ends with 'ith'       |
| Name LIKE 'John'  | Name matches 'John'        |

---

## ... SQL Operators

38/198

Most Unix-based DBMSs utilise the regexp library

- to provide full POSIX regular expression matching

PostgreSQL uses the ~ operator for this:

*Attr* ~ 'RegExp'

PostgreSQL also provides full-text searching (see doc)

---

## ... SQL Operators

39/198

Examples (using POSIX regular expressions):

|                   |                            |
|-------------------|----------------------------|
| Name ~ '^Ja'      | Name begins with 'Ja'      |
| Name ~ '^i'       | Name has 'i' as 2nd letter |
| Name ~ '.*o.*o.*' | Name contains two 'o's     |
| Name ~ 'ith\$'    | Name ends with 'ith'       |
| Name ~ 'John'     | Name matches 'John'        |

---

## ... SQL Operators

40/198

**String manipulation:**

- *str*<sub>1</sub> || *str*<sub>2</sub> ... return concatenation of *str*<sub>1</sub> and *str*<sub>2</sub>
- lower(*str*) ... return lower-case version of *str*
- substring(*str*,*start*,*count*) ... extract chars from *str*

Etc. etc. ... consult your local SQL Manual (e.g. PostgreSQL Sec 9.4)

Note that above operations are null-preserving (strict):

- if any operand is NULL, result is NULL
- beware of (a || ' ' || b || ' ' || c) ... NULL if any of a, b, c are null

---

## ... SQL Operators

41/198

Arithmetic operations:

+ - \* / abs ceil floor power sqrt sin

Aggregations apply to a column of numbers in a relation:

- count(*attr*) ... number of rows in *attr* column
- sum(*attr*) ... sum of values for *attr*
- avg(*attr*) ... mean of values for *attr*
- min/max(*attr*) ... min/max of values for *attr*

Note: count applies to columns of non-numbers as well.

---

## ... SQL Operators

42/198

NULL in arithmetic operation always yields NULL, e.g.

3 + NULL = NULL      1 / NULL = NULL

NULL in aggregations is ignored (treated as unknown), e.g.

```
sum(1,2,3,4,5,6)      = 21
sum(1,2,NULL,4,NULL,6) = 13
avg(1,2,3,4,5)        = 3
avg(NULL,2,NULL,4)     = 3
```

---

## The NULL Value

43/198

Expressions containing NULL generally yield NULL.

However, boolean expressions use three-valued logic:

| <i>a</i> | <i>b</i> | <i>a</i> AND <i>b</i> | <i>a</i> OR <i>b</i> |
|----------|----------|-----------------------|----------------------|
| TRUE     | TRUE     | TRUE                  | TRUE                 |
| TRUE     | FALSE    | FALSE                 | TRUE                 |
| TRUE     | NULL     | NULL                  | TRUE                 |
| FALSE    | FALSE    | FALSE                 | FALSE                |
| FALSE    | NULL     | FALSE                 | NULL                 |
| NULL     | NULL     | NULL                  | NULL                 |

---

### ... The NULL Value

44/198

Important consequence of NULL behaviour ...

These expressions do not work as (might be) expected:

```
x = NULL      x <> NULL
```

Both return NULL regardless of the value of x

Can only test for NULL using:

```
x IS NULL      x IS NOT NULL
```

---

### ... The NULL Value

45/198

Other ways PostgreSQL provides for dealing with NULL:

`coalesce(Val1, Val2, ... Valn)`

- returns first non-null value *Val*<sub>*i*</sub>
- useful for providing a "displayable" value for nulls

`nullif(Val1, Val2)`

- returns null if *Val*<sub>1</sub> is equal to *Val*<sub>2</sub>
- can be used to provide inverse of `coalesce()`

---

## SQL: Schemas

46/198

---

## Relational Data Definition

47/198

In order to give a relational data model, we need to:

- describe tables
- describe attributes that comprise tables
- describe any constraints on the data

A *relation schema* defines an individual table.

A *database schema* is a collection of relation schemas that defines the structure of and constraints on an entire database.

---

## ... Relational Data Definition

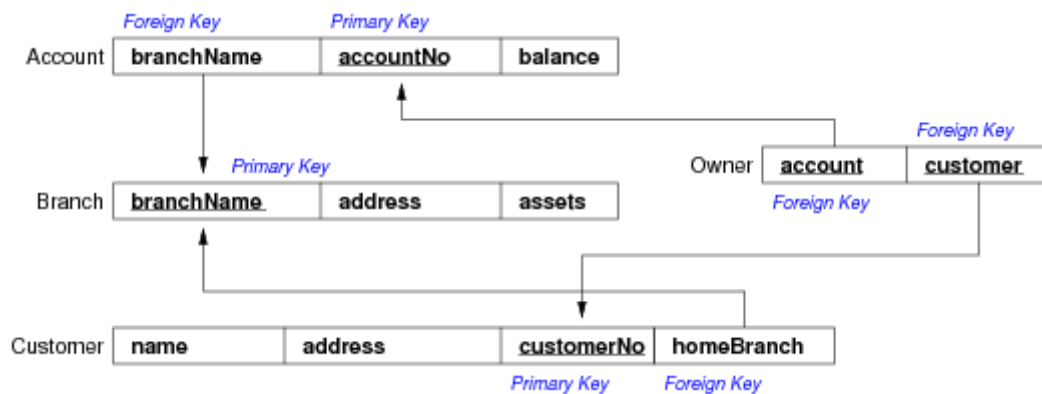
48/198

So far, we have given relational schemas informally, e.g.

- individual relation schemas

```
Account(accountNo, branchName, balance)
Branch(branchNo, address, assets)
Customer(customerNo, name, address, homeBranch)
Owner(customer, branch)
```

- database schemas



---

## SQL Data Definition Language

49/198

SQL is normally considered to be a query language.

However, it also has a data definition sub-language (DDL) for describing database schemas.

The SQL DDL allows us to specify:

- names of tables
- names and domains for attributes
- various types of constraints (e.g. primary/foreign keys)

It also provides mechanisms for performance tuning (see later).

---

## Defining a Database Schema

50/198

Relations (tables) are described using:

```
CREATE TABLE RelName (  
    attribute1 domain1 constraints,  
    attribute2 domain2 constraints,  
    ...  
    table-level constraints, ...  
)
```

where *constraints* can include details about primary keys, foreign keys, default values, and constraints on attribute values.

This not only defines the table schema but also creates an empty instance of the table.

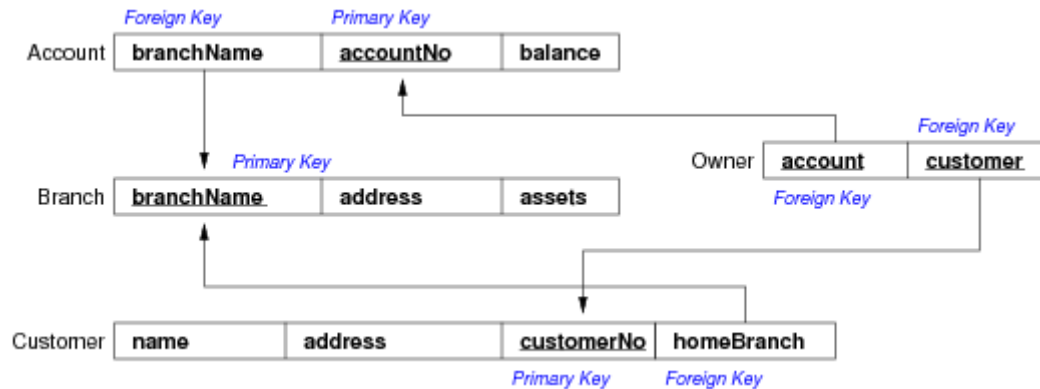
Tables are removed via `DROP TABLE RelName;`

---

## ... Defining a Database Schema

51/198

Consider the relational diagram for the example schema:



This shows explicitly the connection between foreign key attributes and their corresponding key attributes.

The SQL DDL provides notation for expressing this in the table definition.

---

## ... Defining a Database Schema

52/198

SQL DDL for the example schema:

```
CREATE TABLE Branch (  
    name          varchar(30),  
    address       varchar(50),  
    assets        float,  
    PRIMARY KEY   (name)  
);
```

Note: name is required to be unique and not null

---

## ... Defining a Database Schema

53/198

More SQL DDL for the example schema:

```
CREATE TABLE Customer (  
    customerNo    integer,  
    name          varchar(40),  
    address       varchar(50),  
    homeBranch    varchar(30) not null,  
    PRIMARY KEY   (customerNo),  
    FOREIGN KEY   (homeBranch)  
                REFERENCES Branch(name)  
);
```

Note: the not null captures total participation, i.e. every customer has a home branch.

---

## ... Defining a Database Schema

54/198

More SQL DDL for the example schema:

```
CREATE TABLE Account (  
    accountNo     char(5),  
    heldAtBranch  varchar(30) not null,  
    balance       float,  
    PRIMARY KEY   (accountNo),  
    FOREIGN KEY   (heldAtBranch)  
                REFERENCES Branch(name)  
);
```

Note: the not null captures total participation, i.e. every account is held at some branch.

---

More SQL DDL for the example schema:

```
CREATE TABLE OwnedBy (
    account      char(5),
    customer     integer,
    PRIMARY KEY  (account,customer),
    FOREIGN KEY  (account)
                REFERENCES Account(accountNo),
    FOREIGN KEY  (customer)
                REFERENCES Customer(customerNo)
);
```

Note: it is not possible in SQL to capture the semantics that Accounts are required to be owned by some Customer.

## Declaring Keys

56/198

Primary keys:

- if a single attribute, declare with attribute, e.g.  

```
accountNo char(5) PRIMARY KEY,
```
- if several attributes, declare with table constraints, e.g.  

```
name      varchar(40),
address   varchar(50),
...
PRIMARY KEY (name,address)
```

## ... Declaring Keys

57/198

If we want to define a numeric primary key, e.g.

```
CREATE TABLE R ( id INTEGER PRIMARY KEY, ... );
```

we still have the problem of generating unique values.

Most DBMSs provide a mechanism to

- generating a sequence of unique values
- ensuring that tuples don't get assigned the same value

PostgreSQL's version:

```
CREATE TABLE R ( id SERIAL PRIMARY KEY, ... );
```

## ... Declaring Keys

58/198

Foreign keys:

- if a single attribute, specify *Relation(Attribute)*, e.g.  

```
customer integer
            REFERENCES Customer(customerNo)
-- or
customer integer REFERENCES Customer
-- or
FOREIGN KEY (customer)
            REFERENCES Customer(customerNo)
```

## ... Declaring Keys

59/198

Foreign keys: (cont)

- if several attributes, specify in table constraints, e.g.

```

name  varchar(40),
addr  varchar(50),
...
FOREIGN KEY (name,addr)
      REFERENCES Person(name,address)

```

If defining foreign keys with table constraints, must use FOREIGN KEY keywords.

## ... Declaring Keys

60/198

Declaring foreign keys assures **referential integrity**.

Example:

Account.branchName refers to primary key of Branch

If we want to delete a tuple from Branch, and there are tuples in Account that refer to it, we could ...

- **reject** the deletion (PostgreSQL/Oracle default behaviour)
- **set-NULL** the foreign key attributes in Account records
- **cascade** the deletion and remove Account records

## ... Declaring Keys

61/198

Can force the alternative delete behaviours via e.g.

```

-- to cascade deletes
customer integer
      REFERENCES Customer(customerNo)
      ON DELETE CASCADE

```

```

-- to set foreign keys to NULL
customer integer
      REFERENCES Customer(customerNo)
      ON DELETE SET NULL

```

## ... Declaring Keys

62/198

Example of different deletion strategies:

Branch

| branchName            | address              | assets             |
|-----------------------|----------------------|--------------------|
| Downtown              | Brooklyn             | 9000000            |
| Redwood               | Palo Alto            | 2100000            |
| Perryridge            | Horseneck            | 1700000            |
| Mianus                | Horseneck            | 400000             |
| <del>Round Hill</del> | <del>Horseneck</del> | <del>8000000</del> |
| North Town            | Rye                  | 3700000            |
| Brighton              | Brooklyn             | 7100000            |

Account

*After deletion with SET NULL*

| branchName | accountNo | balance |
|------------|-----------|---------|
| Downtown   | A-101     | 500     |
| NULL       | A-215     | 700     |
| Perryridge | A-102     | 400     |
| NULL       | A-305     | 350     |
| Brighton   | A-201     | 900     |
| Redwood    | A-222     | 700     |

Account

*Original relation*

| branchName | accountNo | balance |
|------------|-----------|---------|
| Downtown   | A-101     | 500     |
| Round Hill | A-215     | 700     |
| Perryridge | A-102     | 400     |
| Round Hill | A-305     | 350     |
| Brighton   | A-201     | 900     |
| Redwood    | A-222     | 700     |

Account

*After deletion with CASCADE*

| branchName | accountNo | balance |
|------------|-----------|---------|
| Downtown   | A-101     | 500     |
| Perryridge | A-102     | 400     |
| Brighton   | A-201     | 900     |
| Redwood    | A-222     | 700     |



## Other Attribute Properties

Can specify that an attribute must have a non-null value, e.g.

```
barcode varchar(20) NOT NULL,  
price    float NOT NULL
```

Can specify that an attribute must have a unique value, e.g.

```
barcode varchar(20) UNIQUE,  
isbn     varchar(15) UNIQUE NOT NULL
```

Primary keys are automatically `UNIQUE NOT NULL`.

---

## ... Other Attribute Properties

64/198

Can specify a `DEFAULT` value for an attribute

- will be assigned to attribute if no value is supplied during insert

**Example:**

```
CREATE TABLE Account (  
    accountNo char(5) PRIMARY KEY,  
    branchName varchar(30)  
        REFERENCES Branch(name)  
        DEFAULT 'Central',  
    balance    float DEFAULT 0.0  
);  
  
INSERT INTO Account(accountNo) VALUES ('A-456')  
-- produces the tuple  
Account('A-456', 'Central', 0.0)
```

---

## Attribute Value Constraints

65/198

In fact, `NOT NULL` is a special case of a constraint on the value that an attribute is allowed to take.

SQL has a more general mechanism for specifying such constraints.

```
attrName type CHECK ( condition )
```

The *Condition* can be arbitrarily complex, and may even involve other attributes, relations and `SELECT` queries.

(but many RDBMSs (e.g. Oracle and PostgreSQL) don't allow `SELECT` in `CHECK`)

---

## ... Attribute Value Constraints

66/198

**Example:**

```
CREATE TABLE Example  
(  
    gender CHAR(1) CHECK (gender IN ('M','F')),  
    Xvalue INT    NOT NULL,  
    Yvalue INT    CHECK (Yvalue > Xvalue),  
    Zvalue FLOAT  CHECK (Zvalue >  
                        (SELECT MAX(price)  
                         FROM Sells))  
);
```

---

## Named Constraints

67/198

Any constraint in an SQL DDL can be named via

```
CONSTRAINT constraintName constraint
```

Example:

```
CREATE TABLE Example
(
    gender CHAR(1) CONSTRAINT GenderCheck
                        CHECK (gender IN ('M','F')),
    Xvalue INT          NOT NULL,
    Yvalue INT          CONSTRAINT XYOrder
                        CHECK (Yvalue > Xvalue),
);
```

---

## SQL: Building Databases

68/198

### Creating Databases

69/198

Mechanism for creating databases is typically DBMS-specific.

Many implement a (non-standard) SQL-like statement:

```
CREATE DATABASE DBname;
```

Many provide an external command, e.g PostgreSQL's

```
$ createdb DBname
```

Produces an empty database (no tables, etc) called *DBname*

---

### ... Creating Databases

70/198

A database can be completely removed (no backup) via

```
$ dropdb DBname
```

This removes all tuples, all tables, all traces of *DBname*

Tables can be removed from a database schema via:

```
DROP TableName
```

All tuples can be removed from a table via:

```
DELETE FROM TableName
```

---

### ... Creating Databases

71/198

Loading a schema with PostgreSQL:

```
$ createdb mydb
$ psql mydb
...
mydb=# \i schema.sql
...
```

or

```
$ psql -f schema.sql mydb
```

Running the above as:

```
$ psql -a -f schema.sql mydb
```

intersperses messages with the schema definition.

Useful for debugging, since errors appear in context.

---

## ... Creating Databases

72/198

Re-loading schemas is not well-supported in PostgreSQL.

Simplest approach is:

```
$ dropdb mydb
$ createdb mydb
$ psql -f schema.sql mydb
```

An alternative is to leave DB but drop all tables:

```
$ psql mydb
...
mydb=# drop Table1;
mydb=# drop Table2;
etc. etc. in correct order
mydb=# \i schema.sql
...
```

Later, we'll see how to write functions to automate this.

---

## ... Creating Databases

73/198

The entire contents of a database may be dumped:

```
$ pg_dump mydb > mydb.dump
```

Dumps all definitions needed to re-create entire DB

- table definitions (create table)
- constraints, including PKs and FKs
- all data from all tables
- domains, stored procedures, triggers, etc.

Some things change appearance, but mean the same thing  
(e.g. `varchar(30)` becomes `character varying(30)`, etc.)

---

## ... Creating Databases

74/198

Dumps may be used for backup/restore or copying DBs

```
$ pg_dump mydb > mydb.dump -- backup
$ createdb newdb
$ psql newdb -f mydb.dump -- copy
```

Result: newdb is a snapshot/copy of mydb.

- however, different object identifiers
  - as changes are made, the two DBs will diverge
- 

## Data Modification in SQL

75/198

SQL provides mechanisms for modifying data (tuples) in tables:

- INSERT ... add a new tuple into a table
- DELETE ... remove tuples from a table (via condition)
- UPDATE ... modify values in exiting tuples (via condition)

Constraint checking is applied automatically on any change.

(See description of relational model for details of which checking applied when)

---

Also provides mechanisms for modifying table meta-data:

- CREATE TABLE ... create a new empty table
- DROP TABLE ... remove table from database (incl. tuples)
- ALTER TABLE ... change properties of existing table

Analogous operations are available on other kinds of database objects, e.g.

- CREATE VIEW, CREATE FUNCTION, CREATE RULE, ...
- DROP VIEW, DROP FUNCTION, DROP RULE, ...
- no UPDATE on these; use CREATE OR REPLACE

## Insertion

77/198

Accomplished via the INSERT operation:

```
INSERT INTO RelationName
VALUES (val1, val2, val3, ...)
```

```
INSERT INTO RelationName(Attr1, Attr2, ...)
VALUES (valForAttr1, valForAttr2, ...)
```

Each form adds a single new tuple into *RelationName*.

## ... Insertion

78/198

```
INSERT INTO R VALUES (v1, v2, ...)
```

- values must be supplied for all attributes of *R*
- in same order as appear in CREATE TABLE statement

```
INSERT INTO R(A1, A2, ...) VALUES (v1, v2, ...)
```

- can specify any subset of attributes of *R*
- values must match attribute specification order
- unspecified attributes are assigned default or null

## ... Insertion

79/198

**Example:** Add the fact that Justin likes 'Old'.

```
INSERT INTO Likes VALUES ('Justin', 'Old');
-- or --
INSERT INTO Likes(drinker, beer)
VALUES('Justin', 'Old');
-- or --
INSERT INTO Likes(beer, drinker)
VALUES('Old', 'Justin');
```

**Example:** Add a new drinker with unknown phone number.

```
INSERT INTO Drinkers(name, addr)
VALUES('Frank', 'Coogee');
-- which inserts the tuple ...
('Frank', 'Coogee', null)
```

## ... Insertion

80/198

**Example:** insertion with default values

```
ALTER TABLE Likes
ALTER COLUMN beer SET DEFAULT 'New';
```

```
ALTER TABLE Likes
  ALTER COLUMN drinker SET DEFAULT 'Joe';

INSERT INTO Likes(drinker)
  VALUES('Fred');
INSERT INTO Likes(beer)
  VALUES('Sparkling Ale');

-- inserts the two new tuples ...
('Fred', 'New')
('Joe', 'Sparkling Ale')
```

---

## ... Insertion

81/198

**Example:** insertion with insufficient values.

E.g. specify that drinkers' phone numbers cannot be NULL.

```
ALTER TABLE Drinkers
  ALTER COLUMN phone SET NOT NULL;
```

And then try to insert a new drinker whose phone number we don't know:

```
INSERT INTO Drinkers(name,addr)
  VALUES ('Zoe', 'Manly');
```

```
ERROR: ExecInsert: Fail to add null value
in not null attribute phone
```

---

## Insertion from Queries

82/198

Can use the result of a query to perform insertion of multiple tuples at once.

```
INSERT INTO Relation ( Subquery );
```

Tuples of *Subquery* must be projected into a suitable format (i.e. matching the tuple-type of *Relation* ).

---

## ... Insertion from Queries

83/198

**Example:** Create a relation of potential drinking buddies (i.e. people who go to the same bars as each other).

```
CREATE TABLE DrinkingBuddies (
  drinker varchar(20) references Drinkers(name),
  buddy   varchar(20) references Drinkers(name),
  primary key (drinker,buddy)
);

INSERT INTO DrinkingBuddies (
  SELECT a.drinker AS drinker,
         b.drinker AS buddy
  FROM   Frequents a, Frequents b
  WHERE  a.bar = b.bar AND a.drinker <> b.drinker
);
```

Note: this is better done as a view (treat this as a materialized view).

---

## Bulk Insertion of Data

84/198

Tuples may be inserted individually:

```
insert into Stuff(x,y,s) values (2,4,'green');
insert into Stuff(x,y,s) values (4,8,null);
insert into Stuff(x,y,s) values (8,null,'red');
...
```

but this is tedious if 1000's of tuples are involved.

It is also inefficient, because all relevant constraints are checked after insertion of each tuple.

---

## ... Bulk Insertion of Data

85/198

Most DBMSs provide non-SQL methods for bulk insertion:

- using a compact representation for each tuple
- loading all tuples without constraint checking
- doing all constraint checks at the end

Downside: if even one tuple is buggy, none are inserted.

Example: PostgreSQL's copy statement:

```
copy Stuff(x,y,s) from stdin;
2      4      green
4      8      \N
8      \N      red
\.
```

Can also copy from a named file.

---

## Deletion

86/198

Accomplished via the DELETE operation:

```
DELETE FROM Relation
WHERE Condition
```

Removes all tuples from *Relation* that satisfy *Condition*.

**Example:** Justin no longer likes Sparkling Ale.

```
DELETE FROM Likes
WHERE drinker = 'Justin'
      AND beer = 'Sparkling Ale';
```

**Special case:** Make relation *R* empty.

```
DELETE FROM R;
```

---

## ... Deletion

87/198

**Example:** remove all expensive beers from sale.

```
DELETE FROM Sells
WHERE price >= 3.00;
```

**Example:** remove all drinkers with no fixed address.

```
DELETE FROM Drinkers
WHERE addr IS NULL;
```

This fails if such Drinkers are referenced in other tables.

---

## Semantics of Deletion

88/198

Method A for `DELETE FROM R WHERE Cond` :

```
FOR EACH tuple T in R DO
  IF T satisfies Cond THEN
    remove T from relation R
  END
END
```

Method B for `DELETE FROM R WHERE Cond` :

```

FOR EACH tuple T in R DO
    IF T satisfies Cond THEN
        make a note of this T
    END
END
FOR EACH noted tuple T DO
    remove T from relation R
END

```

Does it matter which method is used?

---

## ... Semantics of Deletion

89/198

**Example:** Delete all beers for which there is another beer by the same manufacturer.

```

DELETE FROM Beers b
WHERE EXISTS
    (SELECT name
     FROM   Beers
     WHERE  manf = b.manf
           AND name <> b.name);

```

Does the query result in ...

- deletion of all beers by brewers who make multiple beers
- deletion of all but the "last beer" by such brewers

Note: PostgreSQL disallows deletions with correlated subqueries (the FROM clause can be only a table name).

---

## ... Semantics of Deletion

90/198

Example continued ...

Different results come from different evaluation methods ..

- Method A: iterate and evaluate condition for each beer
  - consider a manufacturer *M* who makes two beers *A* and *B*
  - when we reach *A*, there are two beers by *M*, so delete *A*
  - when we reach *B*, there are no other beers by *M*, so not deleted
- Method B: evaluate condition and then do all deletions
  - both *A* and *B* test positive, and so both are deleted

Most RDBMSs use Method B, which matches natural semantics of DELETE.

---

## Updates

91/198

An update allows you to modify values of specified attributes in specified tuples of a relation:

```

UPDATE R
SET    list of assignments
WHERE  Condition

```

Each tuple in relation *R* that satisfies *Condition* has the assignments applied to it.

Assignments may:

- assign constant values to attributes,  
e.g. SET price = 2.00
- use existing values in the tuple to compute new values,  
e.g. SET price = price \* 0.5

---

## ... Updates

92/198

**Example:** Adam changes his phone number.

```
UPDATE Drinkers
SET    phone = '9385-2222'
WHERE  name = 'Adam';
```

**Example:** John moves to Coogee.

```
UPDATE Drinkers
SET    addr = 'Coogee',
       phone = '9665-4321'
WHERE  name = 'John';
```

---

## ... Updates

93/198

Can update many tuples at once (all tuples that satisfy condition)

**Example:** Make \$3 the maximum price for beer.

```
UPDATE Sells
SET    price = 3.00
WHERE  price > 3.00;
```

**Example:** Increase beer prices by 10%.

```
UPDATE Sells
SET    price = price * 1.10;
```

---

## Changing Tables

94/198

Accomplished via the ALTER TABLE operation:

ALTER TABLE *Relation Modifications*

Some possible modifications are:

- add a new column (attribute) (set value to NULL unless default given)
- change properties of an existing attribute (e.g. constraints)
- remove an attribute

---

## ... Changing Tables

95/198

**Example:** Add phone numbers for hotels.

```
ALTER TABLE Bars
ADD phone char(10) DEFAULT 'Unlisted';
```

This appends a new column to the table and sets value for this attribute to 'Unlisted' in every tuple.

Specific phone numbers can subsequently be added via:

```
UPDATE Bars
SET    phone = '9665-0000'
WHERE  name = 'Coogee Bay Hotel';
```

If no default value is given, new column is set to all NULL.

---

## For More Details ...

96/198

Full details are in the PostgreSQL Reference Manual.

See the section "SQL Commands", which has entries for

- INSERT, DELETE, UPDATE
- CREATE X, DROP Y, ALTER Z

You will become very familiar with some of these commands by end of session.

---



## Queries

98/198

A *query* is a *declarative program* that retrieves data from a database.

Analogous to an expression in relational algebra.

But SQL does not implement relational algebra precisely.

Queries are used in two ways in RDBMSs:

- interactively (e.g. in `psql`)
  - the entire result is displayed in tabular format on the output
- by a program (e.g. in a PLpgSQL function)
  - the result tuples are consumed one-at-a-time by the program

## Queries in SQL

99/198

The most common kind of SQL statement is the `SELECT` query:

```
SELECT attributes
FROM   relations
WHERE  condition
```

The result of this statement is a relation, which is typically displayed on output.

The `SELECT` statement contains the functionality of select, project and join from the relational algebra.

## SELECT Example

100/198

The question "What beers are made by Toohey's?", can be phrased:

```
SELECT Name FROM Beers WHERE Manf = 'Toohey's';
```

This gives a subset of the Beers relation, displayed as:

```
      name
-----
New
Old
Red
Sheaf Stout
```

Notes:

- upper- and lower-case are not distinguished, except in strings.
- quotes are escaped by doubling them (''' is like C '\')

## Semantics of SELECT

101/198

For SQL `SELECT` statement on a single relation:

```
SELECT Attributes
FROM   R
WHERE  Condition
```

Formal semantics (relational algebra):

$$\text{Proj}[\text{Attributes}](\text{Sel}[\text{Condition}](R))$$

## ... Semantics of SELECT

102/198

Operationally, we think in terms of a *tuple variable* ranging over all tuples of the relation.

Operational semantics:

```
FOR EACH tuple T in R DO
  check whether T satisfies the condition
                        in the WHERE clause
  IF it does THEN
    print the attributes of T that are
      specified in the SELECT clause
  END
END
```

---

## Projection in SQL

103/198

For a relation  $R$  and attributes  $X \subseteq R$ , the relational algebra expression  $\pi_X(R)$  is implemented in SQL as:

```
SELECT X FROM R
```

**Example:** Names of drinkers =  $\pi_{Name}(Drinkers)$

```
SELECT Name FROM Drinkers;
```

```
  name
-----
Adam
Gernot
John
Justin
```

---

## ... Projection in SQL

104/198

**Example:** Names/addresses of drinkers =  $\pi_{Name,Addr}(Drinkers)$

```
SELECT Name, Addr FROM Drinkers;
```

```
  name |   addr
-----+-----
Adam   | Randwick
Gernot | Newtown
John   | Clovelly
Justin | Mosman
```

---

## ... Projection in SQL

105/198

The symbol \* denotes a list of *all* attributes.

**Example:** All information about drinkers =  $(Drinkers)$

```
SELECT * FROM Drinkers;
```

```
  name |   addr   |   phone
-----+-----+-----
Adam   | Randwick | 9385-4444
Gernot | Newtown  | 9415-3378
John   | Clovelly | 9665-1234
Justin | Mosman   | 9845-4321
```

---

## Renaming via AS

106/198

SQL implements renaming () via the AS clause within SELECT.

**Example:** rename Beers(name,manf) to Beers(beer,brewer)

```
SELECT name AS beer, manf AS Brewer
FROM Beers;
```

| beer                | brewer        |
|---------------------|---------------|
| 80/-                | Caledonian    |
| Bigfoot Barley Wine | Sierra Nevada |
| Burraborang Bock    | George IV Inn |
| Crown Lager         | Carlton       |
| Fosters Lager       | Carlton       |
| ...                 |               |

## Expressions as Values in Columns

107/198

AS can also be used to introduce *computed* values (generalised projection)

**Example:** display beer prices in Yen, rather than dollars

```
SELECT bar, beer, price*120 AS PriceInYen FROM Sells;
```

| bar              | beer             | priceinyen       |
|------------------|------------------|------------------|
| Australia Hotel  | Burraborang Bock | 420              |
| Coogee Bay Hotel | New              | 270              |
| Coogee Bay Hotel | Old              | 300              |
| Coogee Bay Hotel | Sparkling Ale    | 335.999994277954 |
| Coogee Bay Hotel | Victoria Bitter  | 275.999994277954 |
| Lord Nelson      | Three Sheets     | 450              |
| Lord Nelson      | Old Admiral      | 450              |
| ...              |                  |                  |

## Text in Result Table

108/198

Trick: to put specific text in output columns

- use string constant expression with AS

**Example:** using Likes(drinker, beer)

```
SELECT drinker, 'likes Cooper's' AS WhoLikes
FROM Likes
WHERE beer = 'Sparkling Ale';
```

| drinker | wholikes       |
|---------|----------------|
| Gernot  | likes Cooper's |
| Justin  | likes Cooper's |

## Selection in SQL

109/198

The relational algebra expression  $\sigma_{Cond}(Rel)$  is implemented in SQL as:

```
SELECT * FROM Rel WHERE Cond
```

**Example:** All about the bars at The Rocks

```
SELECT * FROM Bars WHERE Addr='The Rocks';
```

| name            | addr      | license |
|-----------------|-----------|---------|
| Australia Hotel | The Rocks | 123456  |
| Lord Nelson     | The Rocks | 123888  |
| (2 rows)        |           |         |

The condition can be an arbitrarily complex boolean-valued expression using the operators mentioned previously.

**Example:** Find the price that The Regent charges for New

```
SELECT price
FROM   Sells
WHERE  bar = 'Regent Hotel' AND beer = 'New';

price
-----
2.2
```

This can be formatted better via `to_char`, e.g.

```
SELECT to_char(price, '$99.99') AS price
FROM   Sells
WHERE  bar = 'Regent Hotel' AND beer = 'New';

price
-----
$ 2.20
```

`to_char()` supports a wide range of conversions.

## Multi-relation SELECT Queries

111/198

Syntax is similar to simple SELECT queries:

```
SELECT Attributes
FROM   R1, R2, ...
WHERE  Condition
```

Difference is that FROM clause contains a list of relations.

Also, the condition typically includes cross-relation (join) conditions.

## ... Multi-relation SELECT Queries

112/198

**Example:** Find the brewers whose beers John likes.

```
SELECT Manf as brewer
FROM   Likes, Beers
WHERE  beer = name AND drinker = 'John';
```

```
brewer
-----
Caledonian
Sierra Nevada
Sierra Nevada
Lord Nelson
```

Note: duplicates could be eliminated by using `DISTINCT`.

## ... Multi-relation SELECT Queries

113/198

The above example corresponds to a relational algebra evaluation like:

```
BeerDrinkers = Likes Join[beer=name] Beers
JohnsBeers   = Sel[drinker=John](BeerDrinkers)
Brewers       = Proj[manf](JohnsBeers)
Result       = Rename[manf->brewer](Brewers)
```

The SQL compiler knows how to translate tests

- involving attributes from two relations into a join
- involving attributes from one relations into a selection

---

## Semantics of Multi-Relation SELECT

114/198

For SQL SELECT statement on several relations:

```
SELECT Attributes
FROM   R1, R2, ... Rn
WHERE  Condition
```

Formal semantics (relational algebra):

$$\text{Proj}[\text{Attributes}](\text{Sel}[\text{Condition}](R1 \times R2 \times \dots Rn))$$

---

## ... Semantics of Multi-Relation SELECT

115/198

Operational semantics of SELECT:

```
FOR EACH tuple T1 in R1 DO
  FOR EACH tuple T2 in R2 DO
    ...
    check WHERE condition for current
      assignment of T1, T2, ... vars
    IF holds THEN
      print attributes of T1, T2, ...
        specified in SELECT
    END
  ...
END
```

Requires one tuple variable for each relation, and nested loops over relations. This is *not* how it's actually computed!

---

## Name Clashes in Conditions

116/198

If a selection condition

- refers to two relations
- the relations have attributes with the same name

use the relation name to disambiguate.

**Example:** Which hotels have the same name as a beer?

```
SELECT Bars.name
FROM   Bars, Beers
WHERE  Bars.name = Beers.name;
```

(The answer to this query is empty, but there is nothing special about this)

---

## ... Name Clashes in Conditions

117/198

Can use such qualified names, even if there is no ambiguity:

```
SELECT Sells.beer
FROM   Sells
WHERE  Sells.price > 3.00;
```

Advice:

- qualify attribute names only when absolutely necessary

Note:

- SQL's AS operator is only for renaming output
  - it provides no help with disambiguation
-

The relation-dot-attribute convention doesn't help if we happen to use the same relation twice in a SELECT.

To handle this, we need to define new names for each "instance" of the relation in the FROM clause.

Syntax:

```
SELECT r1.a, r2.b
FROM   R r1, R r2
WHERE  r1.a = r2.a
```

---

### ... Explicit Tuple Variables

119/198

**Example:** Find pairs of beers by the same manufacturer.

```
SELECT b1.name, b2.name
FROM   Beers b1, Beers b2
WHERE  b1.manf = b2.manf AND b1.name < b2.name;
```

| name          | name             |
|---------------|------------------|
| Crown Lager   | Fosters Lager    |
| Crown Lager   | Invalid Stout    |
| Crown Lager   | Melbourne Bitter |
| Crown Lager   | Victoria Bitter  |
| Fosters Lager | Invalid Stout    |
| Fosters Lager | Melbourne Bitter |
| ...           |                  |

The second part of the condition is used to avoid:

- pairing a beer with itself e.g. (New,New)
- same pairs with different order e.g. (New,Old) (Old,New)

---

### ... Explicit Tuple Variables

120/198

A common alternative syntax for

```
SELECT r1.a, r2.b
FROM   R r1, R r2
WHERE  r1.a = r2.a
```

uses the as keyword

```
SELECT r1.a, r2.b
FROM   R as r1, R as r2
WHERE  r1.a = r2.a
```

---

## Explicit Joins

121/198

SQL supports syntax for explicit joins:

```
SELECT...FROM A natural join B
SELECT...FROM A join B using (A1,...,An)
SELECT...FROM A join B on Condition
```

The natural join and join using forms assume that the join attributes are named the same in each relation.

---

### ... Explicit Joins

122/198

**Example:** Find the beers sold at bars where John drinks

```
SELECT Sells.bar, beer, price
FROM   Sells, Frequent
```

```
WHERE drinker = 'John'
      AND Sells.bar = Frequents.bar;
```

could also be expressed as

```
SELECT bar, beer, price
FROM   Sells natural join Frequents
WHERE  drinker='John';
      -- joins on the only common attribute: bar
```

---

## ... Explicit Joins

123/198

The example could also be expressed as

```
SELECT bar, beer, price
FROM   Sells join Frequents using (bar)
WHERE  drinker='John';
      -- only one bar attribute in join result
```

or

```
SELECT Sells.bar, beer, price
FROM   Sells join Frequents
      on Sells.bar = Frequents.bar
WHERE  drinker='John';
      -- bar attribute occurs twice in join result
```

---

## Outer Join

124/198

Join only produces tuples where there are matching values in both of the relations involved in the join.

Often, it is useful to produce results for all tuples in one relation, even if it has no matches in the other.

Consider the query: for each region, find out who drinks there.

---

## ... Outer Join

125/198

A regular join only gives results for regions where people drink.

```
SELECT B.addr, F.drinker
FROM   Bars as B join Frequents as F
      on (bar = name)
ORDER BY addr;
```

| addr      | drinker |
|-----------|---------|
| Coogee    | Adam    |
| Coogee    | John    |
| Kingsford | Justin  |
| Sydney    | Justin  |
| The Rocks | John    |

But what if we want a result that shows all regions, even if there are no drinkers there?

---

## ... Outer Join

126/198

An *outer join* solves this problem.

For  $R$  OUTER JOIN  $S$

- all "tuples" in  $R$  have an entry in the result
- if a tuple from  $R$  matches a tuple in  $S$ , we get the normal join result tuple
- if a tuple from  $R$  has no matches in  $S$ , the attributes supplied by  $S$  are NULL

This outer join variant is called LEFT OUTER JOIN.

---

## ... Outer Join

127/198

Solving the example query with an outer join:

```
SELECT B.addr, F.drinker
FROM   Bars as B
       left outer join
       Frequents as F
       on (bar = name)
ORDER BY B.addr;
```

| addr      | drinker |
|-----------|---------|
| Coogee    | Adam    |
| Coogee    | John    |
| Kingsford | Justin  |
| Randwick  |         |
| Sydney    | Justin  |
| The Rocks | John    |

Note that Randwick is now mentioned (because of the Royal Hotel).

---

## ... Outer Join

128/198

Many RDBMSs provide three variants of outer join:

- *R* LEFT OUTER JOIN *S*
  - behaves as described above
- *R* RIGHT OUTER JOIN *S*
  - includes all tuples from *S* in the result
  - NULL-fills any *S* tuples with no matches in *R*
- *R* FULL OUTER JOIN *S*
  - includes all tuples from *R* and *S* in the result
  - those without matches in other relation are NULL-filled

---

## Subqueries

129/198

The result of a SELECT-FROM-WHERE query can be used in the WHERE clause of another query.

**Simplest Case:** Subquery returns a single, unary tuple

Can treat the result as a single constant value and use in expressions.

Syntax:

```
SELECT *
FROM   R
WHERE  R.a = (SELECT x FROM S WHERE Cond)
        -- assume only one result
```

---

## ... Subqueries

130/198

**Example:** Find bars that serve New at the same price as the Coogee Bay Hotel charges for VB.

```
SELECT bar
FROM   Sells
WHERE  beer = 'New' AND
       price =
       (SELECT price
        FROM   Sells
        WHERE  bar = 'Coogee Bay Hotel'
              AND beer = 'Victoria Bitter');
```

bar



-----  
Royal Hotel

The inner query finds the price of VB at the CBH, and uses this as an argument to a test in the outer query.

---

## ... Subqueries

131/198

Note the potential ambiguity in references to attributes of Sells

```
SELECT bar
FROM   Sells
WHERE  beer = 'New' AND
      price =
        (SELECT price
         FROM   Sells
         WHERE  bar = 'Coogee Bay Hotel'
              AND beer = 'Victoria Bitter');
```

This introduces notions of scope: an attribute refers to the most closely nested relation with that attribute.

Parentheses around the subquery are required (and set the scope).

---

## ... Subqueries

132/198

Note also that the query could be answered via:

```
SELECT s1.bar
FROM   Sells as s1, Sells as s2
WHERE  s1.beer = 'New'
      AND s1.price = s2.price
      AND s2.bar = 'Coogee Bay Hotel'
      AND s2.beer = 'Victoria Bitter';
```

In general, expressing a query via joins will be much more efficient than expressing it with sub-queries.

---

## ... Subqueries

133/198

**Complex Case:** Subquery returns multiple unary tuples.

Treat it as a list of values, and use the various operators on lists/sets (e.g. IN).

**Complex Case:** Subquery returns a relation.

Most of the "list operators" also work on relations.

---

## The IN Operator

134/198

Tests whether a specified tuple is contained in a relation.

*tuple* IN *relation* is true iff the tuple is contained in the relation.

Conversely for *tuple* NOT IN *relation*.

Syntax:

```
SELECT *
FROM   R
WHERE  R.a IN (SELECT x FROM S WHERE Cond)
        -- assume multiple results
```

---

## ... The IN Operator

135/198

**Example:** Find the name and brewer of beers that John likes.

```

SELECT *
FROM   Beers
WHERE  name IN
      (SELECT beer
       FROM   Likes
       WHERE  drinker = 'John');

```

| name                | manf          |
|---------------------|---------------|
| 80/-                | Caledonian    |
| Bigfoot Barley Wine | Sierra Nevada |
| Pale Ale            | Sierra Nevada |
| Three Sheets        | Lord Nelson   |

The subexpression answers the question "What are the names of the beers that John likes?"

## ... The IN Operator

136/198

Note that this query can be answered equally well without using IN.

```

SELECT Beers.name, Beers.manf
FROM   Beers, Likes
WHERE  Likes.drinker = 'John' AND
      Likes.beer = Beers.name;

```

| name                | manf          |
|---------------------|---------------|
| 80/-                | Caledonian    |
| Bigfoot Barley Wine | Sierra Nevada |
| Pale Ale            | Sierra Nevada |
| Three Sheets        | Lord Nelson   |

The version with the subquery corresponds more closely to the way the original query was expressed, and is probably "more natural".

The subquery version is, however, potentially less efficient.

## The EXISTS Function

137/198

EXISTS(*relation*) is true iff the relation is non-empty.

**Example:** Find the beers that are the unique beer by their manufacturer.

```

SELECT name, manf
FROM   Beers b1
WHERE  NOT EXISTS
      (SELECT *
       FROM   Beers
       WHERE  manf = b1.manf
       AND   name != b1.name);

```

Note the scoping rule: to refer to outer Beers in the inner subquery, we need to define a named tuple variable (in this example b1).

A subquery that refers to values from a surrounding query is called a **correlated subquery**.

## Quantifiers

138/198

ANY and ALL behave as existential and universal quantifiers respectively.

**Example:** Find the beers sold for the highest price.

```

SELECT beer
FROM   Sells
WHERE  price >=
      ALL(SELECT price FROM sells);

```

Beware: in common use, "any" and "all" are often synonyms.

E.g. "I'm better than any of you" vs. "I'm better than all of you".

---

## Union, Intersection, Difference

139/198

SQL implements the standard set operations on "union-compatible" relations:

|                   |                                      |
|-------------------|--------------------------------------|
| $R1 \cup R2$      | set of tuples in either $R1$ or $R2$ |
| $R1 \cap R2$      | set of tuples in both $R1$ and $R2$  |
| $R1 \setminus R2$ | set of tuples in $R1$ but not $R2$   |

Oracle deviates from the SQL standard and uses MINUS for EXCEPT; PostgreSQL follows the standard.

---

## ... Union, Intersection, Difference

140/198

**Example:** Find the drinkers and beers such that the drinker likes the beer and frequents a bar that sells it.

```
(SELECT * FROM Likes)
INTERSECT
(SELECT drinker,beer
 FROM Sells natural join Frequents);
```

| drinker | beer            |
|---------|-----------------|
| Adam    | New             |
| John    | Three Sheets    |
| Justin  | Victoria Bitter |

---

## Bag Semantics of SQL

141/198

An SQL relation is really a **bag** (multiset):

- it may contain the same tuple more than once
- unlike lists, there is no specified order on the elements
- example:  $\{1, 2, 1, 3\}$  is a bag and is not a set

This changes the semantics of the "set" operators UNION, INTERSECT and MINUS.

---

## ... Bag Semantics of SQL

142/198

### Bag Union

Sum the times an element appears in the two bags

- example:  $\{1,2,1\} \cup \{1,2,3\} = \{1,1,1,2,2,3\}$

### Bag Intersection

Take the minimum number of occurrences from each bag.

- example:  $\{1,2,1\} \cap \{1,2,3\} = \{1,2\}$

### Bag Difference

Proper-subtract the number of occurrences in the two bags.

- example:  $\{1,2,1\} - \{1,2,3\} = \{1\}$

---

## Forcing Bag/Set Semantics

143/198

Default result for SELECT-FROM-WHERE is a bag.

Default result for UNION, INTERSECT, MINUS is a set.

Why the difference?

A bag can be produced faster because no need to worry about eliminating duplicates (which typically requires sorting).

Can force set semantics with `SELECT DISTINCT`.

Can force bag semantics with `UNION ALL, ...`

---

## ... Forcing Bag/Set Semantics

144/198

**Example:** What beer manufacturers are there?

```
SELECT DISTINCT manf FROM Beers;
```

```
      manf
-----
Caledonian
Carlton
Cascade
Cooper's
George IV Inn
Lord Nelson
Sierra Nevada
Toohey's
```

Note that the result is sorted.

If we omit `DISTINCT`, we get 18 unsorted tuples in the result.

---

## Division

145/198

Not all SQL implementations provide a divide operator, but the same effect can be achieved by combination of existing operations.

**Example:** Find bars that each sell all of the beers Justin likes.

```
SELECT DISTINCT a.bar
FROM   Sells a
WHERE  NOT EXISTS (
    (SELECT beer FROM Likes
     WHERE drinker = 'Justin')
  EXCEPT
  (SELECT beer FROM Sells b
   WHERE bar = a.bar)
);
```

---

## Selection with Aggregation

146/198

Selection clauses can contain aggregation operations.

**Example:** What is the average price of New?

```
SELECT AVG(price)
FROM   Sells
WHERE  beer = 'New';
```

```
      avg
-----
2.38749998807907
```

Note:

- the bag semantics of SQL gives the correct result here

- the price for New in all hotels will be included, even if two hotels sell it at the same price
- if we used set semantics, we'd get the average of all the *different* prices for New.

---

## ... Selection with Aggregation

147/198

If we want set semantics, we can force using DISTINCT.

**Example:** How many different bars sell beer?

```
SELECT COUNT(DISTINCT bar)
FROM Sells;
```

```
count
-----
6
```

Without DISTINCT, the result is 15 ... the number of entries in the Sells table.

---

## Aggregation operators

148/198

The following operators apply to a list (bag) of *numeric* values in one column of a relation:

SUM    AVG    MIN    MAX    COUNT

The notation COUNT(\*) gives the number of tuples in a relation.

**Example:** How many different beers are there?

```
SELECT COUNT(*) FROM Beers;
```

```
count
-----
18
```

---

## Grouping

149/198

SELECT-FROM-WHERE can be followed by GROUP BY to:

- partition result relation into groups  
(according to values of specified attribute)
- summarise some (several) aspects of each group
- output relation contains one tuple per group

**Example:** How many beers does each brewer make?

There is one entry for each beer by each brewer in the Beers table ...

---

## ... Grouping

150/198

The following gives us a list of brewers:

```
SELECT manf FROM Beers;
```

The number of occurrences of each brewer is the number of beers that they make.

Ordering the list makes it much easier to work out:

```
SELECT manf FROM Beers ORDER BY manf;
```

but we still need to count length of runs by hand.

---

## ... Grouping

151/198

If we *group* the runs, we can count(\*) them:

```
SELECT  manf, COUNT(manf)
FROM    Beers
GROUP BY manf;
```

| manf          | count |
|---------------|-------|
| Caledonian    | 1     |
| Carlton       | 5     |
| Cascade       | 1     |
| Cooper's      | 2     |
| George IV Inn | 1     |
| Lord Nelson   | 2     |
| Sierra Nevada | 2     |
| Toohey's      | 4     |

## ... Grouping

152/198

GROUP BY is used as follows:

```
SELECT  attributes/aggregations
FROM    relations
WHERE   condition
GROUP BY attribute
```

Semantics:

1. apply product and selection as for SELECT-FROM-WHERE
2. partition result into groups based on values of *attribute*
3. apply any aggregation separately to each group

## ... Grouping

153/198

The query

```
select manf,count(manf) from Beers group by manf;
```

first produces a partitioned relation and then counts the number of tuples in each partition:

| Name             | Manf       |   | Name                | Manf          |   |
|------------------|------------|---|---------------------|---------------|---|
| 80-              | Caledonian | 1 | Burraborang Bock    | George IV Inn | 1 |
| Crown Lager      | Carlton    | 5 | Old Admiral         | Lord Nelson   | 2 |
| Fosters Lager    | Carlton    |   | Three Sheets        | Lord Nelson   |   |
| Invalid Stout    | Carlton    |   | Bigfoot Barley Wine | Sierra Nevada | 2 |
| Melbourne Bitter | Carlton    |   | Pale Ale            | Sierra Nevada |   |
| Victoria Bitter  | Carlton    |   | New                 | Toohey's      | 4 |
| Premium Lager    | Cascade    | 1 | Old                 | Toohey's      |   |
| Sparkling Ale    | Coopers    | 2 | Red                 | Toohey's      |   |
| Stout            | Coopers    |   | Sheaf Stout         | Toohey's      |   |

## ... Grouping

154/198

Grouping is typically used in queries involving the phrase "for each".

**Example:** For each drinker, find the average price of New at the bars they go to.

```
SELECT  drinker, AVG(price) as "Avg.Price"
FROM    Frequents, Sells
WHERE   beer = 'New'
        AND Frequents.bar = Sells.bar
GROUP BY drinker;
```

| drinker | Avg.Price |
|---------|-----------|
| Adam    | 2.25      |
| John    | 2.25      |
| Justin  | 2.5       |

## Restrictions on SELECT Lists

155/198

When using grouping, every attribute in the SELECT list must:

- have an aggregation operator applied to it OR
- appear in the GROUP-BY clause

**Incorrect Example:** Find the hotel that sells 'New' cheapest.

```
SELECT bar, MIN(price)
FROM   Sells
WHERE  beer = 'New';
```

PostgreSQL's response to this query:

```
ERROR: Attribute sells.bar must be GROUPed
       or used in an aggregate function
```

---

## ... Restrictions on SELECT Lists

156/198

How to answer the query: Which bar sells 'New' cheapest?

```
SELECT bar
FROM   Sells
WHERE  beer = 'New' AND
       price <= (SELECT MIN(price)
                  FROM   Sells
                  WHERE  beer = 'New');
```

```
bar
-----
Regent Hotel
```

---

## ... Restrictions on SELECT Lists

157/198

Also, cannot use grouping to simply re-order results.

**Incorrect Example:** Print beers grouped by their manufacturer.

```
SELECT name, manf FROM Beers
GROUP BY manf;
```

```
ERROR: Attribute beers.name must be GROUPed
       or used in an aggregate function
```

---

## ... Restrictions on SELECT Lists

158/198

How to print beers grouped by their manufacturer?

```
SELECT name, manf FROM Beers
ORDER BY manf;
```

| name             | manf       |
|------------------|------------|
| 80/-             | Caledonian |
| Crown Lager      | Carlton    |
| Fosters Lager    | Carlton    |
| Invalid Stout    | Carlton    |
| Melbourne Bitter | Carlton    |
| Victoria Bitter  | Carlton    |
| Premium Lager    | Cascade    |
| ...              |            |

ORDER BY can be applied to multiple attributes.

---

In some queries, you can use the WHERE condition to eliminate groups.

**Example:** Average beer price by suburb excluding hotels in The Rocks.

```
SELECT Bars.addr, AVG(Sells.price)
FROM Sells, Bars
WHERE Bars.addr != 'The Rocks'
      AND Sells.bar = Bars.name
GROUP BY Bars.addr;
```

For more complex conditions on groups, use the HAVING clause.

## ... Eliminating Groups

160/198

HAVING is used to qualify a GROUP-BY clause:

```
SELECT  attributes/aggregations
FROM    relations
WHERE   condition (on tuples)
GROUP BY attribute
HAVING  condition (on group);
```

Semantics of HAVING:

1. generate the groups as for GROUP-BY
2. eliminate groups *not* satisfying HAVING condition
3. apply aggregations to remaining groups

Note: HAVING condition can use relations/variables from FROM just like WHERE condition, but variables range over each group.

## ... Eliminating Groups

161/198

**Example:** Find the average price of common beers (i.e. those that are served in more than one hotel).

```
SELECT beer,
       to_char(AVG(price), '9.99')
       as "$$$"
FROM   Sells
GROUP BY beer
HAVING COUNT(bar) > 1;
```

| beer            | \$\$\$ |
|-----------------|--------|
| New             | 2.39   |
| Old             | 2.53   |
| Victoria Bitter | 2.40   |

## ... Eliminating Groups

162/198

The HAVING condition can have components that do not use aggregation.

**Example:** Find the average price of beers that are either commonly served (in more than one hotel) or are manufactured by Cooper's.

```
SELECT beer, AVG(price)
FROM   Sells
GROUP BY beer
HAVING COUNT(bar) > 1
      OR beer in
      (SELECT name
       FROM   beers
       WHERE  manf = 'Cooper''s');
```

| beer | avg |
|------|-----|
|------|-----|



|                 |                  |
|-----------------|------------------|
| New             | 2.38749998807907 |
| Old             | 2.53333330154419 |
| Sparkling Ale   | 2.79999995231628 |
| Victoria Bitter | 2.39999997615814 |

## ... Eliminating Groups

163/198

GROUP-BY and HAVING also provide an alternative formulation for division.

**Example:** Find bars that each sell all of the beers Justin likes.

```
SELECT DISTINCT S.bar
FROM   Sells S, Likes L
WHERE  S.beer = L.beer
      AND L.drinker = 'Justin'
GROUP BY bar
HAVING count(S.beer) =
      (SELECT count(beer) FROM Likes
       WHERE drinker = 'Justin');
```

## Partitions and Window Functions

164/198

Sometimes it is useful to

- partition a table into groups
- compute results that apply to each group
- use these results with individual tuples in the group

Comparison with GROUP-BY

- GROUP-BY produces one tuple for each group
- PARTITION augments each tuple with group-based value(s)
- can use other functions than aggregates (e.g. ranking)
- can use attributes other than the partitioning ones

## ... Partitions and Window Functions

165/198

Syntax for PARTITION:

```
SELECT attr1, attr2, ...,
      aggregate1 OVER (PARTITION BY attri),
      aggregate2 OVER (PARTITION BY attrj), ...
FROM   Table
WHERE  condition on attributes
```

Note: the *condition* cannot include the *aggregate* value(s)

## ... Partitions and Window Functions

166/198

Example: show each city with daily temperature and temperature range

Schema: *Weather(city,date,temperature)*

```
SELECT city, date, temperature as temp,
      min(temperature) OVER (PARTITION BY city) as lowest,
      max(temperature) OVER (PARTITION BY city) as highest
FROM   Weather;
```

Output: *Result(city, date, temp, lowest, highest)*

## ... Partitions and Window Functions

167/198

Example showing GROUP BY and PARTITION difference:

```
SELECT city, min(temperature) max(temperature)
FROM Weather GROUP BY city
```

Result: one tuple for each city *Result(city,min,max)*

```
SELECT city, date, temperature as temp,
       min(temperature) OVER (PARTITION BY city),
       max(temperature) OVER (PARTITION BY city)
FROM Weather;
```

Result: one tuple for each temperature measurement.

---

## ... Partitions and Window Functions

168/198

Example: get a list of low-scoring students in each course  
(low-scoring = mark is less than average mark for class)

Schema: *Enrolment(course,student,mark)*

Approach:

- generate tuples containing *(student,mark,classAvg)*
- select just those tuples satisfying *(mark < classAvg)*

Implementation of first step via window function

```
SELECT course, student, mark,
       avg(mark) OVER (PARTITION BY course)
FROM Enrolments;
```

We now look at several ways to complete this data request ...

---

## Complex Queries

169/198

For complex queries, it is often useful to

- break the query into a collection of smaller queries
- define the top-level query in terms of these

This can be accomplished in three ways in SQL:

- views (discussed in detail below)
- subqueries in the FROM clause
- subqueries in a WITH clause

Note that we cannot "correlate" such subqueries in the same way as we can subqueries in the WHERE clause.

---

## ... Complex Queries

170/198

Defining complex queries using views:

```
CREATE VIEW CourseMarksAndAverages(course,student,mark,avg)
AS
SELECT course, student, mark,
       avg(mark) OVER (PARTITION BY course)
FROM Enrolments;
```

```
SELECT course, student, mark
FROM CourseMarksAndAverages
WHERE mark < avg;
```

---

## ... Complex Queries

171/198

In the general case:

```
CREATE VIEW View1(a,b,c,d) AS Query1;  
CREATE VIEW View2(e,f,g) AS Query2;  
...  
SELECT a,f FROM View1, View2 WHERE c = e;
```

Notes:

- look like tables ("virtual" tables)
- exist as objects in the database (stored queries)
- useful if specific query is required frequently

---

## ... Complex Queries

172/198

Defining complex queries using FROM subqueries:

```
SELECT course, student, mark  
FROM (SELECT course, student, mark,  
        avg(mark) OVER (PARTITION BY course)  
        FROM Enrolments) AS CourseMarkAndAverages  
WHERE mark < avg;
```

Avoids the need to define views.

---

## ... Complex Queries

173/198

In the general case:

```
SELECT attributes  
FROM (Query1) AS X,  
      (Query2) AS Y,  
...  
WHERE X.a = Y.b AND other conditions
```

Notes:

- must provide name for each subquery, even if never used
- subquery table inherits attribute names from query  
(e.g. in the above, we assume that Query<sub>1</sub> returns an attribute called a)

---

## ... Complex Queries

174/198

Defining complex queries using WITH:

```
WITH CourseMarksAndAverages AS  
    (SELECT course, student, mark,  
            avg(mark) OVER (PARTITION BY course)  
            FROM CourseEnrolments)  
SELECT course, student, mark, avg  
FROM CourseMarksAndAverages  
WHERE mark < avg;
```

Avoids the need to define views.

---

## ... Complex Queries

175/198

In the general case:

```
WITH Name1(a,b,c) AS (Query1),  
      Name2 AS (Query1), ...  
SELECT attributes
```

```
FROM   Name1, Name2, ...
WHERE  conditions with attributes of Name1 and Name2
```

Notes:

- *Name*<sub>1</sub>, etc. are like temporary tables
- named tables inherit attribute names from query

---

## Recursive Queries

176/198

WITH also provides the basis for recursive queries.

Recursive queries are structured as:

```
WITH RECURSIVE Recurs(attributes) AS (
    SELECT ... not involving Recurs
    UNION
    SELECT ... FROM Recurs, ...
)
SELECT attributes
FROM   Recurs, ...
WHERE  condition involving Recurs attributes
```

Useful for scenarios in which we need to traverse multi-level relationships.

---

### ... Recursive Queries

177/198

Simple example involving a "virtual" table.

Sum the numbers from 1 to 100:

```
WITH RECURSIVE t(n) AS (
    SELECT 1
    UNION
    SELECT n+1 FROM t WHERE n < 100
)
SELECT sum(n) FROM t;
```

---

### ... Recursive Queries

178/198

In the general case:

```
WITH RECURSIVE Recurs(attributes) AS (
    Q1 (non-recursive query)
    UNION
    Q2 (recursive query)
)
SELECT * FROM Recurs;
```

Requires the use of several temporary tables:

- *Result* is the final result of evaluating the query
- *Working*, *Temp* hold intermediate results

---

### ... Recursive Queries

179/198

How recursion works:

```
Working = Result = evaluate Q1
while (Working table is not empty) {
    Temp = evaluate Q2, using Working in place of Recurs
    Temp = Temp - Result
    Result = Result UNION Temp
```

```
    Working = Temp
}
```

I.e. generate new tuples until we see nothing not already seen.

---

## ... Recursive Queries

180/198

Example: count number of each sub-part in a given part.

Schema: *Parts(part, sub\_part, quantity)*

```
WITH RECURSIVE IncludedParts(sub_part, part, quantity) AS (
    SELECT sub_part, part, quantity
    FROM   Parts WHERE part = GivenPart
    UNION ALL
    SELECT p.sub_part, p.part, p.quantity
    FROM   IncludedParts i, Parts p
    WHERE  p.part = i.sub_part
)
SELECT sub_part, SUM(quantity) as total_quantity
FROM   IncludedParts
GROUP BY sub_part
```

---

## SQL: Views

181/198

### Views

182/198

A *view* is like a "virtual relation" defined via a query.

View definition and removal:

```
CREATE VIEW ViewName AS Query
```

```
CREATE VIEW ViewName [ (AttributeNames) ]
AS Query
```

```
DROP VIEW ViewName
```

The *Query* may be any SQL query, involving

- other views (*intensional relations*)
  - stored tables (*extensional relations*)
- 

## ... Views

183/198

The stored tables in a view are referred to as *base tables*.

Views are defined only after their base tables are defined.

A view is valid only as long as its underlying query is valid.

Dropping a view has no effect on the base tables.

---

## ... Views

184/198

**Example:** An avid Carlton drinker might not be interested in any other kinds of beer.

```
CREATE VIEW MyBeers AS
    SELECT name, manf
    FROM   Beers
    WHERE  manf = 'Carlton';

SELECT * FROM MyBeers;
```

| name             | manf    |
|------------------|---------|
| Crown Lager      | Carlton |
| Fosters Lager    | Carlton |
| Invalid Stout    | Carlton |
| Melbourne Bitter | Carlton |
| Victoria Bitter  | Carlton |

## ... Views

185/198

A view might not use all attributes of the base relations.

**Example:** We don't really need the address of inner-city hotels.

```
CREATE VIEW InnerCityHotels AS
  SELECT name, license
  FROM   Bars
  WHERE  addr in ('The Rocks','Sydney');
```

```
SELECT * FROM InnerCityHotels;
```

| name            | license |
|-----------------|---------|
| Australia Hotel | 123456  |
| Lord Nelson     | 123888  |
| Marble Bar      | 122123  |

## ... Views

186/198

A view might use computed attribute values.

**Example:** Number of beers produced by each brewer.

```
CREATE VIEW BeersBrewed AS
  SELECT manf as brewer,
         count(*) as nbeers
  FROM   beers GROUP BY manf;
```

```
SELECT * FROM BeersBrewed;
```

| brewer     | nbeers |
|------------|--------|
| Caledonian | 1      |
| Carlton    | 5      |
| Cascade    | 1      |
| ...        |        |

## Renaming View Attributes

187/198

This can be achieved in two different ways:

```
CREATE VIEW InnerCityHotels AS
  SELECT name AS pub, license AS lic
  FROM   Bars
  WHERE  addr IN ('The Rocks', 'Sydney');
```

```
CREATE VIEW InnerCityHotels(pub,lic) AS
  SELECT name, license
  FROM   Bars
  WHERE  addr IN ('The Rocks', 'Sydney');
```

## Using Views

188/198

Views can be used in queries as if they were stored relations.

However, they differ from stored relations in two important respects:

- their "value" can change without being explicitly modified  
(i.e. a view may change whenever one of its base tables is updated)
- they may not be able to be explicitly modified (updated)  
(only a certain simple kinds of views can be explicitly updated)

---

## ... Using Views

189/198

**Example:** of view changing when base table changes.

```
SELECT * FROM InnerCityHotels;
      name      | license
-----+-----
Australia Hotel | 123456
Lord Nelson     | 123888
Marble Bar      | 122123
```

```
-- then the Lord Nelson goes broke
DELETE FROM Bars WHERE name = 'Lord Nelson';

-- no explicit update has been made to InnerCityHotels
SELECT * FROM InnerCityHotels;
      name      | license
-----+-----
Australia Hotel | 123456
Marble Bar      | 122123
```

---

## Updating Views

190/198

Explicit updates are allowed on views satisfying the following:

- the view involves a single relation R
- the WHERE clause does not involve R in a subquery
- the WHERE clause only uses attributes from the SELECT

Attributes not in the view's SELECT will be set to NULL in the base relation after an insert into the view.

---

## ... Updating Views

191/198

**Example:** Our InnerCityHotel view is not updatable.

```
INSERT INTO InnerCityHotels
VALUES ('Jackson's on George', '9876543');
```

creates a new tuple in the Bars relation:

```
(Jackson's on George, NULL, 9876543)
```

when we SELECT from the view, this new tuple does not satisfy the view condition:

```
addr IN ('The Rocks', 'Sydney')
```

---

## ... Updating Views

192/198

If we had chosen to omit the license attribute instead, it would be updatable:

```
CREATE VIEW CityHotels AS
  SELECT name,addr FROM Bars
  WHERE addr IN ('The Rocks', 'Sydney');
```

```
INSERT INTO CityHotels
VALUES ('Jackson's on George', 'Sydney');
```

creates a new tuple in the Bars relation:

(Jackson's on George, Sydney, NULL)

which would appear in the view after the insertion.

---

## ... Updating Views

193/198

Updatable views in PostgreSQL require us to specify explicitly how updates are done:

```
CREATE RULE InsertCityHotel AS
  ON INSERT TO CityHotels
  DO INSTEAD
    INSERT INTO Bars VALUES
      (new.name, new.addr, NULL);
```

```
CREATE RULE UpdateCityHotel AS
  ON UPDATE TO CityHotels
  DO INSTEAD
    UPDATE Bars
    SET   addr = new.addr
    WHERE name = old.name;
```

---

## Evaluating Views

194/198

Two alternative ways of implementing views:

- re-writing rules (or macros)
  - when a view is used in a query, the query is re-written
  - after rewriting, becomes a query only on base relations
- explicit stored relations (called *materialized views*)
  - the view is stored as a real table in the database
  - updated appropriately when base tables are modified

The difference: underlying query evaluated either at query time or at update time.

---

## ... Evaluating Views

195/198

**Example:** Using the InnerCityHotels view.

```
CREATE VIEW InnerCityHotels AS
  SELECT name, license
  FROM   Bars
  WHERE  addr IN ('The Rocks', 'Sydney');
```

```
SELECT name
FROM   InnerCityHotels
WHERE  license = '123456';
```

--is rewritten into the following form before execution

```
SELECT name
FROM   Bars
WHERE  addr IN ('The Rocks', 'Sydney')
      AND license = '123456';
```

---

## ... Evaluating Views

196/198

Demonstrate the rewriting process via relational algebra.

Some abbreviations

- $n$  = name,  $l$  = license
- $L$  = license = '123456'
- $A$  = addr IN ('The Rocks', 'Sydney')

View definition in RA:



InnerCityHotels =  $\pi_{(n,L)}(\sigma_{(A)}(\text{Bars}))$

---

## ... Evaluating Views

197/198

Rewriting of query involving a view:

= SELECT name from InnerCityHotels  
WHERE license = '123456'

=  $\pi_{(n)}(\sigma_{(L)}(\text{InnerCityHotels}))$

=  $\pi_{(n)}(\sigma_{(L)}(\pi_{(n,L)}(\sigma_{(A)}(\text{Bars}))))$

=  $\pi_{(n)}(\pi_{(n,L)}(\sigma_{(L)}(\sigma_{(A)}(\text{Bars}))))$

=  $\pi_{(n)}(\sigma_{(L)}(\sigma_{(A)}(\text{Bars})))$

=  $\pi_{(n)}(\sigma_{(L \ \& \ A)}(\text{Bars}))$

=  $\pi_{(n)}(\sigma_{(A \ \& \ L)}(\text{Bars}))$

= SELECT name FROM Bars  
WHERE addr IN ('The Rocks', 'Sydney')  
AND license = '123456'

---

## Materialized Views

198/198

Naive implementation of materialized views:

- replace view table by re-evaluating query after each update

Clearly this costs space and makes updates more expensive.

However, in a situation where

- updates are infrequent compared to queries on the view
- the cost of "computing" the view is expensive

this approach provides substantial benefits.

Materialized views are used extensively in data warehouses.

---