

# COMP3311 Week 06 Lecture

---

## Triggers

---

### Triggers (review)

2/31

*Triggers* are actions invoked by DB modifications.

They allow programmers to

- implement global constraint (assertion) checking
- maintain summary values (cross-table dependencies)

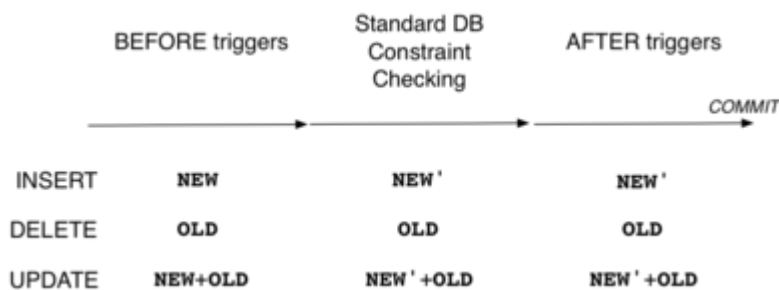
They achieve this by

- invoking functions before/after insert/delete/update
  - using/manipulating OLD/NEW values of changed tuples
- 

### ... Triggers (review)

3/31

Sequence of activities during database update:



Note: BEFORE trigger can modify value of new tuple

---

## Triggers in PostgreSQL

4/31

PostgreSQL triggers provide a mechanism for

- INSERT, DELETE or UPDATE events
- to automatically activate PLpgSQL functions

Syntax for PostgreSQL trigger definition:

```
CREATE TRIGGER TriggerName
{AFTER|BEFORE} Event1 [OR Event2 ...]
ON TableName
[ WHEN ( Condition ) ]
```

```
FOR EACH {ROW|STATEMENT}  
EXECUTE PROCEDURE FunctionName(args...);
```

---

### ... Triggers in PostgreSQL

5/31

There is no restriction on what code can go in the function.

However a BEFORE function must contain one of:

```
RETURN old;      or      RETURN new;
```

depending on which version of the tuple is to be used.

If BEFORE trigger returns old, no change occurs.

If exception is raised in trigger function, no change occurs.

---

### Trigger Example #1

6/31

Consider a database of people in the USA:

```
create table Person (  
    id          integer primary key,  
    ssn         varchar(11) unique,  
    ... e.g. family, given, street, town ...  
    state       char(2), ...  
);  
create table States (  
    id          integer primary key,  
    code        char(2) unique,  
    ... e.g. name, area, population, flag ...  
);
```

Constraint:  $\text{Person.state} \in (\text{select code from States})$ , or  
 $\text{exists (select id from States where code=Person.state)}$

---

### ... Trigger Example #1

7/31

**Example:** ensure that only valid state codes are used:

```
create trigger checkState before insert or update  
on Person for each row execute procedure checkState();  
  
create function checkState() returns trigger as $$  
begin  
    -- normalise the user-supplied value  
    new.state = upper(trim(new.state));
```

```

    if (new.state !~ '^[A-Z][A-Z]$') then
        raise exception 'Code must be two alpha chars';
    end if;
    -- implement referential integrity check
    select * from States where code=new.state;
    if (not found) then
        raise exception 'Invalid code %',new.state;
    end if;
    return new;
end;
$$ language plpgsql;

```

---

### ... Trigger Example #1

8/31

Examples of how this trigger would behave:

```

insert into Person
    values('John',..., 'Calif.',...);
-- fails with 'Statecode must be two alpha chars'

insert into Person
    values('Jane',..., 'NY',...);
-- insert succeeds; Jane lives in New York

update Person
    set town='Sunnyvale', state='CA'
    where name='Dave';
-- update succeeds; Dave moves to California

update Person
    set state='0Z' where name='Pete';
-- fails with 'Invalid state code 0Z'

```

---

### Trigger Example #2

9/31

**Example:** department salary totals

Scenario:

```

Employee(id, name, address, dept, salary, ...)
Department(id, name, manager, totSal, ...)

```

An assertion that we wish to maintain:

```

create assertion TotalSalary check (
    not exists (
        select d.id from Department d
        where d.totSal <>

```

```
        (select sum(e.salary)
         from Employee e
         where e.dept = d.id)
    )
)
```

---

### ... Trigger Example #2

10/31

Events that might affect the validity of the database

- a new employee starts work in some department
- an employee gets a rise in salary
- an employee changes from one department to another
- an employee leaves the company

A single assertion could check for this after each change.

With triggers, we have to program each case separately.

Each program implements updates to *ensure* assertion holds.

---

### ... Trigger Example #2

11/31

Implement the Employee update triggers from above in PostgreSQL:

Case 1: new employees arrive

```
create trigger TotalSalary1
after insert on Employees
for each row execute procedure totalSalary1();

create function totalSalary1() returns trigger
as $$
begin
    if (new.dept is not null) then
        update Department
        set     totSal = totSal + new.salary
        where  Department.id = new.dept;
    end if;
    return new;
end;
$$ language plpgsql;
```

---

### ... Trigger Example #2

12/31

Case 2: employees change departments/salaries

```
create trigger TotalSalary2
```

```
after update on Employee
for each row execute procedure totalSalary2();

create function totalSalary2() returns trigger
as $$
begin
    update Department
    set    totSal = totSal + new.salary
    where Department.id = new.dept;
    update Department
    set    totSal = totSal - old.salary
    where Department.id = old.dept;
    return new;
end;
$$ language plpgsql;
```

---

### ... Trigger Example #2

13/31

Case 3: employees leave

```
create trigger TotalSalary3
after delete on Employee
for each row execute procedure totalSalary3();

create function totalSalary3() returns trigger
as $$
begin
    if (old.dept is not null) then
        update Department
        set    totSal = totSal - old.salary
        where Department.id = old.dept;
    end if;
    return old;
end;
$$ language plpgsql;
```

---

### Exercise: Triggers (1)

14/31

Requirement: maintain assets in bank branches

- each branch has assets based on the accounts held there
- whenever an account changes, the assets of the corresponding branch should be updated to reflect this change

Some possible changes:

- a new account is opened
- the amount of money in an account changes
- an account moves from one branch to another

- an account is closed

Implement triggers to maintain `Branch.assets`  
[\[Solutions\]](#)

---

## Exercise: Triggers (2)

15/31

Consider a simple airline flights/bookings database:

```
Airports(id, code, name, city)
Planes(id, craft, nseats)
Flights(id, fltNum, plane, source, dest
        departs, arrives, price, seatsAvail)
Passengers(id, name, address, phone)
Bookings(pax, flight, paid)
```

Write triggers to ensure that `Flights.seatsAvail` is consistent with number of Bookings on that flight.

Assume that we never UPDATE a booking (only insert/delete)

[\[Solutions\]](#)

---

## Programming with DBs

---

### Programming with Databases

17/31

So far, we have seen ...

- accessing data via SQL queries
- packaging SQL queries as views/functions
- building functions to return tables
- implementing assertions via triggers

All of the above programming

- is very close to the data
  - takes place inside the DBMS
- 

### ... Programming with Databases

18/31

Complete applications require code outside the DBMS

- to handle the user interface (GUI or Web)
- to interact with other systems (e.g. other DBs)
- to perform compute-intensive work (vs. data-intensive)

"Conventional" programming languages (PLs) provide these.

---

## ... Programming with Databases

19/31

Requirements of an interface between PL and RDBMS:

- mechanism for connecting to the DBMS
- mapping between tuples and PL objects
- mechanism for mapping PL "requests" to queries
- mechanism for iterating over query results

Distance between PL and DBMS is variable, e.g.

- libpq allows C programs to use PG structs
  - JDBC transmits SQL strings, retrieves tuples-as-objects
- 

## PL/DB Interface

20/31

Common DB access API used in programming languages

```
db = connect_to_dbms(DBname, User/Password);
query = build_SQL("SqlStatementTemplate", values);
results = execute_query(db, query);
while (more_tuples_in(results))
{
    tuple = fetch_row_from(results);
    // do something with values in tuple ...
}
```

This pattern is used in many different libraries:

- Java/JDBC, PHP/PDO, Perl/DBI, Python/dbapi2, Tcl, ...
- 

## ... PL/DB Interface

21/31

DB access libraries have similar overall structure.

However, they differ in the details:

- whether specific to one database or generic
- whether object-oriented or procedural flavour
- function/method names and parameters
- how to get data from program into SQL statements
- how to get data from tuples to program variables

We use PHP to illustrate the idea in this lecture.

---

Standard pattern for extracting data from DB:

```
$db = dbConnect("dbname=myDB");
...
$query = "select a,b,c from R where c >= %d";
$result = dbQuery($db, mkSQL($query, $min));
while ($tuple = dbNext($result)) {
    $tmp = $tuple["a"] - $tuple["b"] - $tuple["c"];
    # or ...
    list($a,$b,$c) = $tuple;
    $tmp = $a - $b - $c;
}
...
```

---

## DB Library

23/31

Functions in the database library:

- `dbConnect(conn)`: establish connection to DB
- `dbQuery(db,sql)`: send SQL statement for execution
- `dbNext(res)`: fetch next tuple from result set
- `dbUpdate(db,sql)`: send SQL insert/delete/update
- ...

Most functions terminate with message if error occurs.

---

## ... DB Library

24/31

```
$t = dbNext(resource $r);
```

- `$t` is assigned next tuple from result set `$r`
- `$t` contains two copies of values from tuple
  - one set of values is indexed by position in SELECT clause
  - one set of values is indexed by name in SELECT clause

Example:

```
$q = "select name,max(mark) from Enrolments ...";
$r = dbQuery($db,$q);
$t = dbNext($r);
# results in $t with value
array(0=>'John', "name"=>'John', 1=>95, "max"=>95)
```

---

## Example PHP code (actual code)

25/31

```
$db_handle = pg_connect("dbname=bpsimple");
```



```

$query = "SELECT title, fname, lname FROM customer";
$result = pg_exec($db_handle, $query);
if ($result) {
    echo "The query executed successfully.\n";
    for ($row = 0; $row < pg_numrows($result); $row++) {
        $fullname = pg_result($result, $row, 'title') . " ";
        $fullname .= pg_result($result, $row, 'fname') . " ";
        $fullname .= pg_result($result, $row, 'lname');
        echo "Customer: $fullname\n";
    }
} else {
    echo "The query failed with the following error:\n";
    echo pg_errormessage($db_handle);
}
pg_close($db_handle);

```

---

## DB/PL Mismatch

26/31

There is a tension between PLs and DBMSs

- DBMSs deal efficiently with sets of tuples
- PLs encourage dealing with single tuples/objects

If not handled carefully, can lead to inefficient use of DB.

Note: relative costs of DB access operations:

- establishing a DBMS connection ... very high
  - initiating an SQL query ... high
  - accessing individual tuple ... low
- 

## ... DB/PL Mismatch

27/31

Consider the PL/DBMS access method, phrased in PHP:

```

-- establish connection to DBMS
$db = dbAccess("DB");
$query = "select a,b from R,S where ... ";
-- invoke query and get handle to result set
$results = dbQuery($db, $query);
-- for each tuple in result set
while ($tuple = dbNext($results)) {
    -- process next tuple
    process($tuple['a'], $tuple['b']);
}

```

---

## ... DB/PL Mismatch

28/31

Example: find mature-age students

```
$query = "select * from Student";
$results = dbQuery($db,$query);
while ($tuple = dbNext($results)) {
    if ($tuple["age"] >= 40) {
        -- process mature-age student
    }
}
```

If 10000 students, and only 500 of them are over 40, we transfer 9500 unnecessary tuples from DB.

---

### ... DB/PL Mismatch

29/31

E.g. should be implemented as:

```
$query = "select * from Student where age >= 40";
$results = dbQuery($db,$query);
while ($tuple = dbNext($results)) {
    -- process mature-age student
}
```

Transfers only the 500 tuples that are needed.

---

### ... DB/PL Mismatch

30/31

Example: find info about all marks for all students

```
$query1 = "select id,name from Student";
$res1 = dbQuery($db,$query1);
while ($tuple1 = dbNext($res1)) {
    $query2 = "select course,mark from Marks".
              " where student = $tuple1['id']";
    $res2 = dbQuery($db,$query2);
    while ($tuple2 = dbNext($res2)) {
        -- process student/course/mark info
    }
}
```

If 10000 students, we invoke 10001 queries on the database.

---

### ... DB/PL Mismatch

31/31

E.g. should be implemented as:

```
$query = "select id,name,course,mark".
         " from Student s, Marks m".
```

```
        " where s.id = m.student";  
$results = dbQuery($db,$query);  
while ($tuple = dbNext($results)) {  
    -- process student/course/mark info  
}
```

We invoke 1 query, and transfer same number of tuples.

---

Produced: 10 April 2018