

# COMP3311 Week 11 Lecture

---

## Transaction Processing

---

### Transactions, Concurrency, Recovery

2/27

DBMSs provide access to valuable information resources in an environment that is:

- *shared* - concurrent access by multiple users
- *unstable* - potential for hardware/software failure

Each user should see the system as:

- *unshared* - their work is not inadvertently affected by others
- *stable* - the data survives in the face of system failures

Ultimate goal: data integrity is maintained at all times.

---

### ... Transactions, Concurrency, Recovery

3/27

Transaction processing

- techniques for managing "logical units of work" which may require multiple DB operations

Concurrency control

- techniques for ensuring that multiple concurrent transactions do not interfere with each other

Recovery mechanisms

- techniques to restore information to a consistent state, even after major hardware shutdowns/failures
- 

### Transactions

4/27

A *transaction* is

- an atomic "unit of work" in an application
- which may require multiple database changes

Transactions happen in a multi-user, unreliable environment.

To maintain integrity of data, transactions must be:

- Atomic - either fully completed or totally rolled-back
- Consistent - map DB between consistent states
- Isolated - transactions do not interfere with each other
- Durable - persistent, restorable after system failures

---

## Example Transaction

5/27

Bank funds transfer

- move  $N$  dollars from account  $X$  to account  $Y$
- Accounts(id,name,balance,heldAt, ...)
- Branches(id,name,address,assets, ...)
- maintain Branches.assets as sum of balances via triggers
- transfer implemented by function which
  - has three parameters: amount, source acct, dest acct
  - checks validity of supplied accounts
  - checks sufficient available funds
  - returns a unique transaction ID on success

---

## ... Example Transaction

6/27

```
create or replace function
  transfer(N integer, X text, Y text) returns integer
declare
  xID integer; yID integer; avail integer;
begin
  select id,balance into xID,avail
  from Accounts where name=X;
  if (xID is null) then
    raise exception 'Invalid source account %',X;
  end if;
  select id into yID
  from Accounts where name=Y;
  if (yID is null) then
    raise exception 'Invalid dest account %',Y;
  end if;
  ...
```

---

## ... Example Transaction

7/27

```
...
  if (avail < N) then
    raise exception 'Insufficient funds in %',X;
  end if;
  -- total funds in system = NNNN
```

```
update Accounts set balance = balance-N
where id = xID;
-- funds temporarily "lost" from system
update Accounts set balance = balance+N
where id = yID;
-- funds restored to system; total funds = NNNN
return nextval('tx_id_seq');
end;
```


---

## Transaction Concepts

8/27

A transaction must always terminate, either:

- successfully (COMMIT), with all changes preserved
- unsuccessfully (ABORT), with database unchanged



---

## ... Transaction Concepts

9/27

To describe transaction effects, we consider:

- READ - transfer data from disk to memory
- WRITE - transfer data from memory to disk
- ABORT - terminate transaction, unsuccessfully
- COMMIT - terminate transaction, successfully

Normally abbreviated to R(X), W(X), A, C

SELECT produces READ operations on the database.

INSERT produces WRITE operations.

UPDATE, DELETE produce both READ + WRITE operations.


---

## Transaction Consistency

10/27

Transactions typically have intermediate states that are inconsistent.

However, states *before* and *after* transaction must be consistent.



Reminder: "consistent" = satisfying all of the specified constraints

## ... Transaction Consistency

11/27

Transaction descriptions can be abstracted

- consider only *Read* and *Write* operations on shared data
- e.g. T1: R(X) W(X) R(Y) W(Y), T2: R(X) R(Y) W(X) W(Y)

A *schedule* defines

- a specific execution of one or more transactions
- typically concurrent, with interleaved operations

Arbitrary interleaving of operations causes *anomalies*

- two consistency-preserving transactions
- produce a final state which is not consistent

## Serial Schedules

12/27

*Serial* execution: T1 then T2 or T2 then T1

T1: R(X) W(X) R(Y) W(Y)

T2: R(X) W(X)

or

T1: R(X) W(X) R(Y) W(Y)

T2: R(X) W(X)

Serial execution guarantees a consistent final state if

- the initial state of the database is consistent
- T1 and T2 are consistency-preserving

## Concurrent Schedules

13/27

*Concurrent* schedules interleave T1,T2,... operations

Some concurrent schedules are ok, e.g.

T1: R(X) W(X) R(Y) W(Y)

T2:                      R(X)                      W(X)

Other concurrent schedules cause anomalies, e.g.

T1: R(X)                      W(X)                      R(Y) W(Y)  
T2:                      R(X)                      W(X)

Want the system to ensure that only valid schedules occur.

## Serializability

14/27

*Serializable* schedule:

- concurrent schedule for  $T_1 .. T_n$  with final state  $S$
- $S$  is also a final state of one of the possible serial schedules for  $T_1 .. T_n$

Abstracting this needs a notion of *schedule equivalence*.

Two common formulations of *serializability*:

- *conflict serializability* (read/write operations occur in the "right" order)
- *view serializability* (read operations see the correct version of data)

## Conflict Serializability

15/27

Consider two transactions  $T_1$  and  $T_2$  acting on data item  $X$ .

Possible orders for read/write operations by  $T_1$  and  $T_2$ :

<b><math>T_1</math> first</b>	<b><math>T_2</math> first</b>	Equiv?
$R_1(X) R_2(X)$	$R_2(X) R_1(X)$	yes
$R_1(X) W_2(X)$	$W_2(X) R_1(X)$	no
$W_1(X) R_2(X)$	$R_2(X) W_1(X)$	no
$W_1(X) W_2(X)$	$W_2(X) W_1(X)$	no

If  $T_1$  and  $T_2$  act on different data items, result is always equivalent.

## ... Conflict Serializability

16/27

Two transactions have a potential *conflict* if

- they perform operations on the same data item
- at least one of the operations is a write operation

In such cases, the order of operations affects the result.

If no conflict, can swap order without affecting the result.

If we can transform a schedule

- by swapping the order of non-conflicting operations
- such that the result is a serial schedule

then we say that the schedule is *conflict serializable*.

---

### ... Conflict Serializability

17/27

Example: transform a concurrent schedule to serial schedule

```
T1: R(A) W(A)      R(B)      W(B)
T2:      R(A)      W(A)      R(B) W(B)
swap
T1: R(A) W(A) R(B)      W(B)
T2:      R(A) W(A)      R(B) W(B)
swap
T1: R(A) W(A) R(B)      W(B)
T2:      R(A)      W(A) R(B) W(B)
swap
T1: R(A) W(A) R(B) W(B)
T2:      R(A) W(A) R(B) W(B)
```

---

### ... Conflict Serializability

18/27

Checking for conflict-serializability:

- show that ordering in concurrent schedule
- cannot be achieved in any serial schedule

Method for doing this:

- build a *precedence-graph*
- nodes represent transactions
- arcs represent order of action on shared data
- arc from  $T_1 \rightarrow T_2$  means  $T_1$  acts on  $X$  before  $T_2$
- cycles indicate *not* conflict-serializable.

---

## Concurrency Control

---

### Concurrency Control

20/27

Serializability tests are useful theoretically ...

But don't provide a mechanism for organising schedules

- they can only be done "after the event"
- they are computationally very expensive  $O(n!)$

What is required are methods that ...

- can be applied to each transaction individually
- guarantee that overall schedule is serializable

---

## ... Concurrency Control

21/27

Approaches to ensuring ACID transactions:

- lock-based

Synchronise transaction execution via locks on some portion of the database.

- version-based

Allow multiple consistent versions of the data to exist, and allow each transaction exclusive access to one version.

- timestamp-based

Organise transaction execution in advance by assigning timestamps to operations.

- validation-based (optimistic concurrency control)

Exploit typical execution-sequence properties of transactions to determine safety dynamically.

---

## Lock-based Concurrency Control

22/27

Synchronise access to shared data items via following rules:

- before reading  $X$ , get shared (read) lock on  $X$
- before writing  $X$ , get exclusive (write) lock on  $X$
- an attempt to get a shared lock on  $X$  is blocked if another transaction already has exclusive lock on  $X$
- an attempt to get an exclusive lock on  $X$  is blocked if another transaction has any kind of lock on  $X$

These rules alone do not guarantee serializability.

Locking also introduces potential for deadlock and starvation.

---

## Locking and Performance

23/27

Locking reduces concurrency  $\Rightarrow$  lower throughput.

*Granularity* of locking can impact performance:

- + lock a small item  $\Rightarrow$  more of database accessible
- + lock a small item  $\Rightarrow$  quick update  $\Rightarrow$  quick lock release
- lock small items  $\Rightarrow$  more locks  $\Rightarrow$  more lock management

Granularity levels: field, row (tuple), table, whole database

Many DBMSs support multiple lock-granularities.

---

## Multi-version Concurrency Control

24/27

One approach to reducing the requirement for locks is to

- provide multiple (consistent) versions of the database
- give each transaction access to an "appropriate" version (i.e. a version that maintains the serializability of the transaction)

This approach is called *Multi-Version Concurrency Control*.

Differences between MVCC and standard locking models:

- writing never blocks reading (make new version of tuple)
  - reading never blocks writing (read old version of tuple)
- 

## Concurrency Control in SQL

25/27

Transactions in SQL are specified by

- **BEGIN** ... start a transaction
- **COMMIT** ... successfully complete a transaction
- **ROLLBACK** ... undo changes made by transaction + abort

In PostgreSQL, other actions that cause rollback:

- **raise exception** during execution of a function
  - returning null from a **before** trigger
- 

## ... Concurrency Control in SQL

26/27

Concurrent access can be controlled via SQL:

- table-level locking: apply lock to entire table



- row-level locking: apply lock to just some rows

**LOCK TABLE** explicitly acquires lock on an entire table.

Other SQL commands implicitly acquire appropriate locks, e.g.

- ALTER TABLE acquires an exclusive lock on table
- UPDATE, DELETE acquire locks on affected rows

All locks are released at end of transaction (no explicit unlock)

---

## Examples

Schedule Q1




T1:	R(A)	W(A)		R(B)		W(B)
T2:			R(A)		W(A)	R(B) W(B)

Schedule Q2

T1:	R(A)		W(A)		R(B)		W(B)
T2:		R(A)		W(A)		R(B)	W(B)

Schedule Q3

T3:		R(B)				W(B)	
T4:	R(A)		W(A)	R(B)			W(B)
T5:					R(A)	W(A)	



[\[Detailed Solutions\]](#)