

Data Structures II

COMP4128 Programming Challenges

School of Computer Science and Engineering
UNSW Australia

Data
Structures II

Range Queries
and Updates

Lowest
Common
Ancestor

Binary
function
composition

1 Range Queries and Updates

2 Lowest Common Ancestor

3 Binary function composition

- Given N integers a_0, a_1, \dots, a_{N-1} , answer queries of the form:

$$\sum_{i=l}^{r-1} a_i$$

for given pairs l, r .

- N is up to 100,000.
- There are up to 100,000 queries.
- We can't answer each query naïvely, we need to do some kind of precomputation.

- **Algorithm** Construct an array of prefix sums.
- $b_0 = a_0$.
- $b_i = b_{i-1} + a_i$.
- This takes $O(N)$ time.
- Now, we can answer every query in $O(1)$ time.
- This works on any “reversible” operation. That is, any operation $A \star B$ where if we know $A \star B$ and A , we can find B .
- This includes addition and multiplication, but *not* max or gcd.


- We can now receive updates mixed in with the queries. The updates are of the form

$$a_i := k$$

for given i and k .


- There can still be up to 100,000 updates and queries in total. Recomputing the prefix sums will take $O(N)$ time per update, so our previous solution is now $O(N^2)$ for this problem, which is too slow.
- As we have done in the past, we try to find a solution that slows down our queries but speeds up updates in order to improve the overall complexity.

- So far, we've seen trees where each node contains a value we care about. This tree is a little different:




- We only store our data in the leaves. The rest of the tree is intermediate computations.
- Each parent stores the sum of its children.
- Now what? Trees are great, but how do we perform our updates and queries?


- Let's update the element at index 2 to 25.




- Let's update the element at index 2 to 25.




- Let's update the element at index 2 to 25.




- Let's update the element at index 2 to 25.



- Let's update the element at index 2 to 25.




- Let's update the element at index 2 to 25.




- We always construct the tree so that it's balanced, then its height is $O(\log N)$.
- Thus, updates take $O(\log N)$ time.
- Still, all of this is useless if we can't actually query the tree fast. How do we do that?

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- Each node in the tree has a “range of responsibility”.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- Each node in the tree has a “range of responsibility”, split evenly between its children.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- We start at the top of the tree, and 'push' the query range down into the applicable nodes.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- This is a recursive call, so we do one branch at a time. Let's start with the left branch.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- There is no need to continue further into the left subtree, because it doesn't intersect the query range.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- There is also no need to continue further down, because this range is equal to our query range.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- We return the result we have obtained up to the chain, and let the query continue.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- We return the result we have obtained up to the chain, and let the query continue.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- Now, it is time to recurse into the other branch of this query.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- Here, the query range is equal to the node's range of responsibility, so we're done.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).




- Here, the query range is equal to the node's range of responsibility, so we're done.

- Let's query the sum of $[2, 8)$ (inclusive-exclusive).



- Now that we've obtained both results, we can add them together and return the answer.

- We didn't visit many nodes during our query.



- In fact, because only the left and right edges of the query can ever get as far as the leaves, and ranges in the middle stop much higher, we only visit $O(\log N)$ nodes during a query.

- Thus we have $O(\log N)$ time for both updates and queries.
- This data structure is commonly known as a range tree, segment tree, interval tree, tournament tree, etc.
- The number of nodes we add halves on each level, so the total number of nodes is still $O(N)$.
- For ease of understanding, the illustrations used a *full* binary tree, which always has a number of nodes one less than a power-of-two. This data structure works fine as a *complete* binary tree as well (all layers except the last are filled). This case is harder to imagine conceptually but the implementation works fine.
- All this means is that padding out the data to the nearest power of two is not necessary.

- Since these binary trees are complete, they can be implemented using the same array-based tree representation as with an array heap
 - Place the root at index 0. Then for each node i , its children (if they exist) are $2i + 1$ and $2i + 2$.
 - Alternatively, place the root at index 1, then for each node i the children are $2i$ and $2i + 1$.
- This works with any binary associative operator, e.g.
 - min, max
 - sum
 - gcd
 - merge (from merge sort)
 - For a non-constant-time operation like this one, multiply the complexity of all range tree operations by the complexity of the merging operation.

- We can extend range trees to allow range updates in $O(\log N)$ using *lazy propagation*
- The basic idea is similar to range queries: push the update down recursively into the nodes whose range of responsibility intersects the update range.
- However, to keep our $O(\log N)$ time complexity, we can't actually update every value in the range.
- Just like we returned early from queries when the query range matched a node's entire range, we cache the update at that node and return without actually applying it.
- Whenever a query or a subsequent update is performed which visits this node, just push the cached update one level further down.
- This is fiddly to implement as updates pile up and often need to combine. Implementation left as an exercise :)

• Implementation (updates)

```

#define MAX_N 100000
// the number of additional nodes created can be as high as the next
// power of two up from MAX_N (131,072)
int tree[266666];

// a is the index in the array. 0- or 1-based doesn't matter here, as
// long as it is nonnegative and less than MAX_N.
// v is the value the a-th element will be updated to.
// i is the index in the tree, rooted at 1 so children are 2i and 2i
// +1.
// instead of storing each node's range of responsibility, we
// calculate it on the way down.
// the root node is responsible for [0, MAX_N)
void update(int a, int v, int i = 1, int start = 0, int end = MAX_N) {
    // this node's range of responsibility is 1, so it is a leaf
    if (end - start == 1) {
        tree[i] = v;
        return;
    }
    // figure out which child is responsible for the index (a) being
    // updated
    int mid = (start + end) / 2;
    if (a < mid) update(a, v, i * 2, start, mid);
    else update(a, v, i * 2 + 1, mid, end);
    // once we have updated the correct child, recalculate our stored
    // value.
    tree[i] = tree[i*2] + tree[i*2+1];
}

```

- Implementation (queries)

```
// query the sum in [a, b)
int query(int a, int b, int i = 1, int start = 1, int end = MAX_N) {
    // the query range exactly matches this node's range of
    // responsibility
    if (start == a && end == b) return tree[i];
    // we might need to query one or both of the children
    int mid = (start + end) / 2;
    int answer = 0;
    // the left child can query [a, mid)
    if (a < mid) answer += query(a, min(b, mid), i * 2, start, mid);
    // the right child can query [mid, b)
    if (b > mid) answer += query(max(a, mid), b, i * 2 + 1, mid, end);
    return answer;
}
```

- Implementation (construction)

- It is possible to construct a range tree in $O(N)$ time, but anything you use it for will take $O(N \log N)$ time anyway.
- Instead of extra code to construct the tree, just call update repeatedly for $O(N \log N)$ construction.

- Problem statement** Gilderoy Lockhart has fallen upon hard times, and now finds himself performing magic tricks to entertain Muggles. In his newest trick, he places n cards face down on a table, and turns to face away from the table. He then invites q members of the audience to do either of the following moves:
 - announce two numbers i and j , and flip all cards between i and j inclusive, or
 - ask him whether a particular card i is face up or face down.

True to form, Lockhart is unable to do this trick himself, so write a program to help him!
- Input** The numbers n and q , each up to 100,000, followed by q lines either of the form $F \ i \ j$ ($1 \leq i \leq j \leq n$), a flip, or $Q \ i$ ($1 \leq i \leq n$), a query.
- Output** For each query, print “Face up” or “Face down”.

- Observe that we can just keep track of how many times each card was flipped; the parity of this number determines whether it is face up or face down.
- We need to handle flips in faster than linear time.
- If we handle flips by adding 1 at the left endpoint and subtracting 1 at the right endpoint, then the sum up to card i (the prefix sum) is the number of times that card i has been flipped.

- **Algorithm** Construct a range tree. For the move $F\ i\ j$, increment a_i and decrement a_{j+1} , and for the move $Q\ i$, calculate $b_i = a_1 + a_2 + \dots + a_i$ modulo 2.
- **Complexity** Each of these operations takes $O(\log n)$ time, so the time complexity is $O(q \log n)$.

• Implementation

```
#include <iostream>
using namespace std;

int main() {
    int n, q;
    for (int i = 0; i < q; i++) {
        char type;
        cin >> type;
        if (type == 'F') {
            int i, j;
            cin >> i >> j;
            update(i, 1);
            update(j + 1, -1);
        }
        else if (type == 'Q') {
            int i;
            cin >> i;
            printf("%s\n", (query(1, i) % 2) ? "Face up" : "Face down");
        }
    }
}
```

Data
Structures II

Range Queries
and Updates

Lowest
Common
Ancestor

Binary
function
composition

1 Range Queries and Updates

2 Lowest Common Ancestor

3 Binary function composition

- **Problem statement** You are given a labelled rooted tree, T , and Q queries of the form, “What is the vertex furthest away from the root in the tree that is an ancestor of vertices labelled u and v ?”
- **Input** A rooted tree T ($1 \leq |T| \leq 1,000,000$), as well as Q ($1 \leq Q \leq 1,000,000$) pairs of integers u and v .
- **Output** A single integer for each query, the label for the vertex that is furthest away from the root that is an ancestor of u and v

- **Algorithm 1** The most straightforward algorithm to solve this problem involves starting with pointers to the vertices u and v , and then moving them upwards towards the root until they're both at the same depth in the tree, and then moving them together until they reach the same place
- This is $O(n)$ per query, since it's possible we need to traverse the entire height of the tree, which is not bounded by anything useful

- The first step we can take is to try to make the “move towards root” step faster
- Since the tree doesn’t change, we can pre-process the tree somehow so we can jump quickly

Data
Structures II

Range Queries
and Updates

Lowest
Common
Ancestor

Binary
function
composition

- 1 Range Queries and Updates
- 2 Lowest Common Ancestor
- 3 Binary function composition

- Let's examine the parent relation $\text{parent}[u]$ in the tree
- Our “move towards root” operation is really just repeated application of this parent relation
- The vertex two steps above u is $\text{parent}[\text{parent}[u]]$, and three steps above is $\text{parent}[\text{parent}[\text{parent}[u]]]$

- Immediately, we can precompute the values $\text{parent}[u][k]$, which is $\text{parent}[u]$ applied k times
- This doesn't have an easy straightforward application to our problem, nor is it fast enough for our purposes

- If we only precompute $\text{parent}[u][k]$ for each $k = 2^\ell$, we only need to perform $O(\log n)$ computations.
- Then, we can then compose up to $\log n$ of these precomputed values to obtain $\text{parent}[u][k]$ for arbitrary k
- To see this, write out the binary expansion of k and keep greedily striking out the most significant set bit — there are at most $\log n$ of them.

- **Algorithm 2** Instead of walking up single edges, we use our precomputed $\text{parent}[u][k]$ to keep greedily moving up by the largest power of 2 possible until we're at the desired vertex
- We need $O(n \log n)$ time and memory to preprocess the required compositions of our parent relation in the tree, as well as $O(\log n)$ time to handle each query

• Implementation (preprocessing)

```
// parent[u][k] is the 2^k-th parent of u
void preprocess() {
    for (int i = 0; i < n; i++) {
        // assume parent[i][0] (the parent of i) is already filled in
        for (int j = 1; (1<<j) < n; j++) {
            parent[i][j] = -1;
        }
    }

    // fill in the parent for each power of two up to n
    for (int j = 1; (1<<j) < n; j++) {
        for (int i = 0; i < n; i++) {
            if (parent[i][j-1] != -1) {
                // the 2^j-th parent is the 2^(j-1)-th parent of the 2^(j-1)-th parent
                parent[i][j] = parent[parent[i][j-1]][j-1];
            }
        }
    }
}
```

• Implementation (querying)

```
int lca (int u, int v) {
    // make sure u is deeper than v
    if (depth[u] < depth[v]) swap(u,v);

    // log[i] holds the largest k such that 2^k <= i
    for (int i = log[depth[u]]; i >= 0; i--) {
        // repeatedly raise u by the largest possible power of two until
        // it is the same depth as v
        if (depth[u] - (1<<i) >= depth[v]) u = parent[u][i];
    }

    if (u == v) return u;

    for (i = log[depth[u]]; i >= 0; i--)
        if (parent[u][i] != -1 && parent[u][i] != parent[v][i]) {
            // raise u and v as much as possible without having them
            // coincide
            // this is important because we're looking for the lowest common
            // ancestor, not just any
            u = parent[u][i];
            v = parent[v][i];
        }

    // u and v are now distinct but have the same parent, and that
    // parent is the LCA
    return parent[u][0];
}
```